

Data Preparation

Paul Rodriguez



On the Importance of Data Prep

- **“Garbage in, garbage out”**
- **Sometimes takes 60-80% of the whole data mining effort**

Working definition

- **Data Preparation:**
 - Cleaning
 - Filtering
 - Transforming
 - Organizing the data matrix (aka 'data wrangling' or 'data munging')

In a nutshell, preparing data for modeling

Cleaning Noise

- Entity Resolution and Record Linkage

e.g. Are these equal?

West Main Street


W Main St

Strategy:

use dictionaries and search possible matches

Statistical Noise:

- Outliers
e.g. remove them,




mean + 3*std-devm

Statistical Noise:

- Outliers

e.g. remove them, but cutoff is arbitrary



mean + 3*std-dev

Missing Data

- **Important to review statistics of a missing variable**
 - Are entries missing completely at random?
 - Are entries missing depending on some other variable?
 - What are the counts and combinations of missing entries among variables?
 - Is there a relation between missing cases and outcome variable?

Missing Data

- **Not applicable**

e.g. spouse name depends on marital status

- **Not Available**

unknown

not entered

Missing Data

- Do missing cases depend on some other variable?
e.g. 'CEOs' don't like to list their salary

Strategy: *get most common job titles
for missing salaries*

Quick Approaches

- Delete instances
and/or
- Delete attributes with high missingness

Quick Approaches

- Leave as 'NULL' category
 - Some algorithms implementation handle NULL (ie Decision Trees)

Simple Imputation

- Use the attribute mean (by class)

Complicated Imputation

- Use a model (based on other attributes) to infer missing value

Complicated Imputation

- Use a model (based on other attributes) to infer missing value

*Best strategy depends on
time vs accuracy tradeoffs*

R and missing data

- **Several packages, such as ‘mice’ , ‘amelia’**
- **Produces multiple data sets**
- **Iterates over missing data estimates and linear model estimates**
 - Mice uses Gibbs sampling (slower)
 - Amelia uses Expectation Maximization (faster)
- **Beware of correlation in variables**
 - Matrices not invertible

R and missing data

- **Sample R code using Amelia:**

Data: UN conflict data in pairs of countries

300K rows ~ 1 hour on Gordon compute node (not run on the user's PC)

1K-100K entries missing per col for about 20 of 50 cols

Note: mice package is probably more well known, and has similar options, but MCMC is slower than EM

```
# run the imputation
library('amelia')
a.out <- amelia(data, ts = "year", cs = "dyadid",
               idvars = c("dyadidyr", "cntryera", "statea", "stateb"),
               intercs=FALSE, p2s = 2, m=10, parallel = "multicore")
```

time variable → `ts = "year"`

crossection → `cs = "dyadid"`



ignore 'id' variables → `idvars = c("dyadidyr", "cntryera", "statea", "stateb")`

interactions → `intercs=FALSE`

parallel options → `p2s = 2, m=10, parallel = "multicore"`

R and missing data

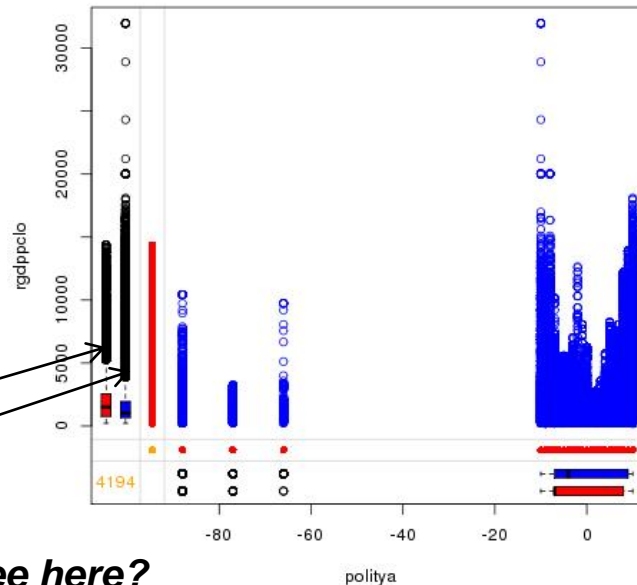
#QA on missing data by comparing density of imputed & original data
`compare.density(a.out, var="politya")`
`compare.density(a.out, var='rgdppcontg')`



R and missing data

```
# Useful library for printing margin plots, to compare histograms  
# conditional on missing/non-missing data  
library('VIM')  
marginplot(gart2use[,c('politya', 'rgdppclo')],  
           col=c('blue','red','orange'))
```

*dist of rgdppclo when
politya missing
politya present*



Variable Transformations

- **Engineer new features**
- **Combine attributes**
e.g. rates and ratios
- **Normalize or Scale data**
- **Discretize data**
(perhaps more intuitive to deal with binned data)

Feature Engineering is Variable Enhancement

- Use Domain and world knowledge

- Examples:

given date and location of doctor visits

- deduce a new variable for Number-of-1st-time-visits
- deduce a new variable for Number-of-visits-over-25-miles
- deduce a new variable for Amount-of-time-between-visits

Re-scaling

- **Mean center** $x_{new} = x - \text{mean}(x)$
- **z-score** $score = \frac{x - \text{mean}(x)}{\text{std}(x)}$
- **Scale to [0...1]** $x_{new} = \frac{x - \min(x)}{\max(x) - \min(x)}$
- **log scaling** $x_{new} = \log(x)$

Variable selection

- **Heuristic methods:**
 - remove variables with low correlations to outcome
(other criteria: information gain, sensitivity, etc...)
- step wise: add 1 variable at a time and test algorithm on samples

Variable selection

- **Some algorithms are robust to extra noise variables**
- **E.g. Least Angle Regression (L_1 penalty),**
 - penalize small effect sizes (zero them out)
- **E.g. Random Forest outputs ‘importance’**
 - low importance implies non-influence in the model
(other criteria: information gain, sensitivity, etc...)

Summary

- **Preparing data is based on statistical principles,**
- **But also heuristics**

Data Prep Exercise: Weather Data

weather - Excel

FILE HOME INSERT PAGE LAYOUT FORMULAS DATA REVIEW VIEW ACROBAT

Clipboard Font Alignment Number Styles Cells Editing

AutoSum Fill Clear Sort & Find & Filter Select

A1

	A	B	C	D	E	F	G	H	I	J
1	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDi	WindGustSp	WindDir9am
2	11/1/2007	Canberra	8	24.3	0	3.4	6.3	NW	30	SW
3	11/2/2007	Canberra	14	26.9	3.6	4.4	9.7	ENE	39	E
4	11/3/2007	Canberra	13.7	23.4	3.6	5.8	3.3	NW	85	N
5	11/4/2007	Canberra	13.3	15.5	39.8	7.2	9.1	NW	54	WNW
6	11/5/2007	Canberra	7.6	16.1	2.8	5.6	10.6	SSE	50	SSE
7	11/6/2007	Canberra	6.2	16.9	0	5.8	8.2	SE	44	SE
8	11/7/2007	Canberra	6.1	18.2	0.2	4.2	8.4	SE	43	SE
9	11/8/2007	Canberra	8.3	17	0	5.6	4.6	E	41	SE
10	11/9/2007	Canberra	8.8	19.5	0	4	4.1	S	48	E
11	11/10/2007	Canberra	8.4	22.8	16.2	5.4	7.7	E	31	S
12	11/11/2007	Canberra	9.1	25.2	0	4.2	11.9	N	30	SE
13	11/12/2007	Canberra	8.5	27.3	0.2	7.2	12.5	E	41	E
14	11/13/2007	Canberra	10.1	27.9	0	7.2	13	WNW	30	S

weather

READY AVERAGE: 2330.477965 COUNT: 8808 SUM: 15325223.1 100%

Search the web and Windows

10:16 AM 8/3/2016

Data Prep exercise

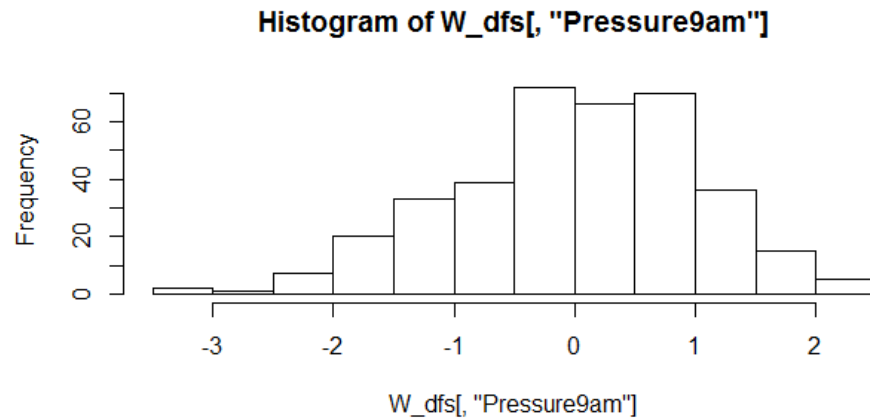

```
W_df = read.table('weather.csv',  
                  header=TRUE,  
                  sep="," ,  
                  stringsAsFactors = TRUE)
```

#Try: run the summary command, what do you notice about value ranges?

```
#Let do some normalization
#Make a function
myscale = function(x)
  if (class(x)=='integer' || class(x)=='numeric') {
    (x-mean(x,na.rm=T)) / sd(x,na.rm=T) }
  else { x}

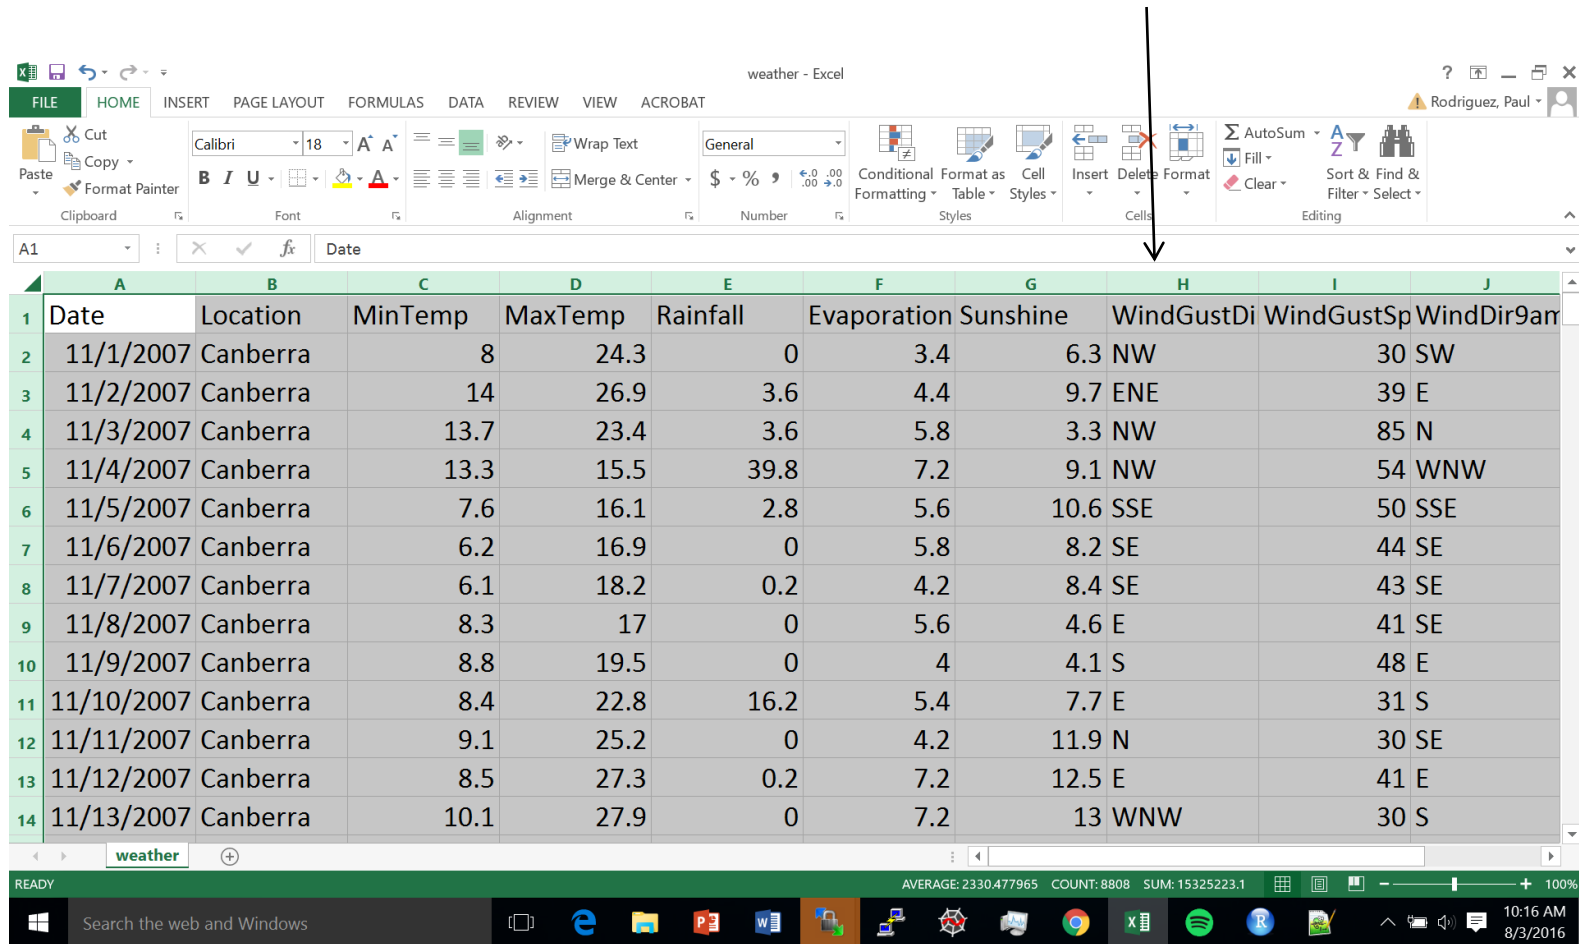
#get a new dataframe and replace with normalized
#  values
W_dfs = W_df
for (i in num_classes)
  { W_dfs[,i]=myscale(W_dfs[,i])}
```

```
hist(W_df[, 'Pressure9am'])
```



Transforming Weather Data Matrix

Let's consider WindGustDir as if it's a repeated measurement

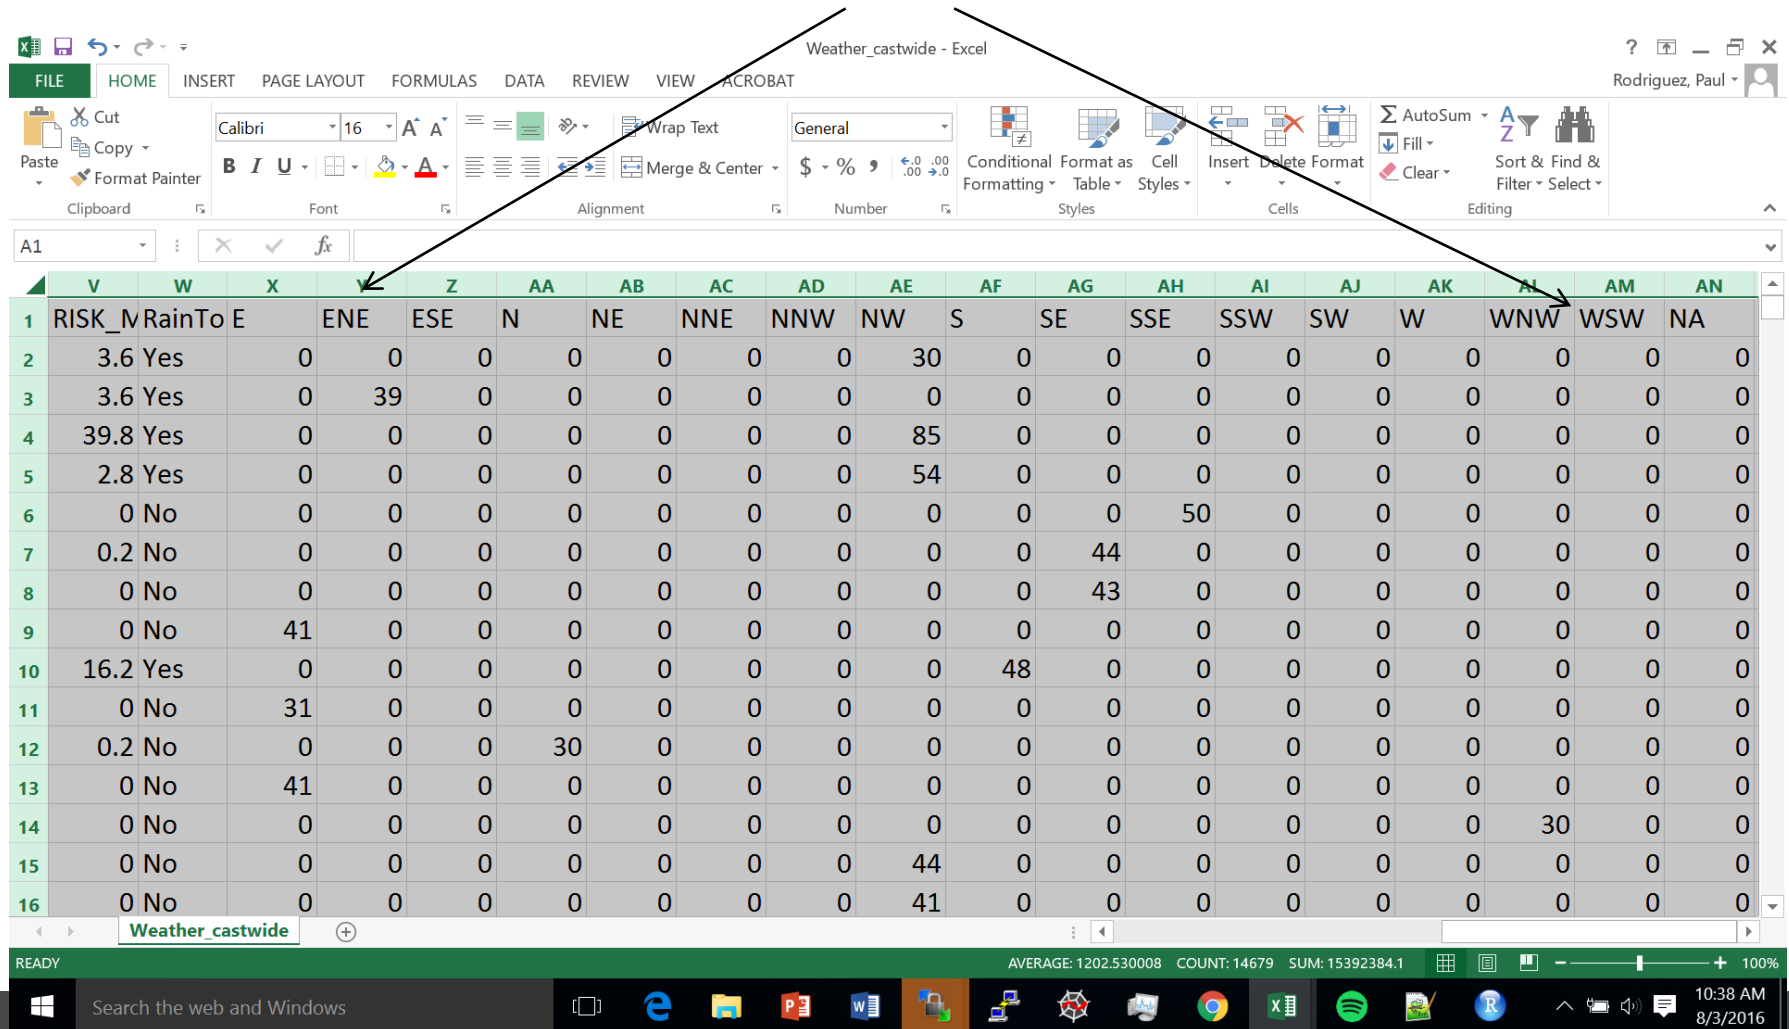


The screenshot shows an Excel spreadsheet with the following data:

	A	B	C	D	E	F	G	H	I	J
1	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSp	WindDir9am
2	11/1/2007	Canberra	8	24.3	0	3.4	6.3	NW	30	SW
3	11/2/2007	Canberra	14	26.9	3.6	4.4	9.7	ENE	39	E
4	11/3/2007	Canberra	13.7	23.4	3.6	5.8	3.3	NW	85	N
5	11/4/2007	Canberra	13.3	15.5	39.8	7.2	9.1	NW	54	WNW
6	11/5/2007	Canberra	7.6	16.1	2.8	5.6	10.6	SSE	50	SSE
7	11/6/2007	Canberra	6.2	16.9	0	5.8	8.2	SE	44	SE
8	11/7/2007	Canberra	6.1	18.2	0.2	4.2	8.4	SE	43	SE
9	11/8/2007	Canberra	8.3	17	0	5.6	4.6	E	41	SE
10	11/9/2007	Canberra	8.8	19.5	0	4	4.1	S	48	E
11	11/10/2007	Canberra	8.4	22.8	16.2	5.4	7.7	E	31	S
12	11/11/2007	Canberra	9.1	25.2	0	4.2	11.9	N	30	SE
13	11/12/2007	Canberra	8.5	27.3	0.2	7.2	12.5	E	41	E
14	11/13/2007	Canberra	10.1	27.9	0	7.2	13	WNW	30	S

Transforming Weather Data Matrix

Now: WindGustDir values each have their own column



Weather_castwide - Excel

FILE HOME INSERT PAGE LAYOUT FORMULAS DATA REVIEW VIEW MACROBOTS

Clipboard Font Alignment Number Styles Cells Editing

Calibri 16 A A

General

Conditional Formatting Format as Table Cell Styles

Insert Delete Format

AutoSum Fill Clear

Sort & Find & Filter Select

A1

	V	W	X	Y	Z	AA	AB	AC	AD	AE	AF	AG	AH	AI	AJ	AK	AL	AM	AN
1	RISK_N	RainTo	E	ENE	ESE	N	NE	NNE	NNW	NW	S	SE	SSE	SSW	SW	W	WNW	WSW	NA
2	3.6	Yes		0	0	0	0	0	0	0	30	0	0	0	0	0	0	0	0
3	3.6	Yes		0	39	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	39.8	Yes		0	0	0	0	0	0	85	0	0	0	0	0	0	0	0	0
5	2.8	Yes		0	0	0	0	0	0	54	0	0	0	0	0	0	0	0	0
6	0	No		0	0	0	0	0	0	0	0	0	50	0	0	0	0	0	0
7	0.2	No		0	0	0	0	0	0	0	0	44	0	0	0	0	0	0	0
8	0	No		0	0	0	0	0	0	0	0	43	0	0	0	0	0	0	0
9	0	No	41	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
10	16.2	Yes		0	0	0	0	0	0	0	48	0	0	0	0	0	0	0	0
11	0	No	31	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
12	0.2	No		0	0	30	0	0	0	0	0	0	0	0	0	0	0	0	0
13	0	No	41	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
14	0	No		0	0	0	0	0	0	0	0	0	0	0	0	0	30	0	0
15	0	No		0	0	0	0	0	0	44	0	0	0	0	0	0	0	0	0
16	0	No		0	0	0	0	0	0	41	0	0	0	0	0	0	0	0	0

Weather_castwide

READY AVERAGE: 1202.530008 COUNT: 14679 SUM: 15392384.1 100%

Search the web and Windows

SDSC SAN DIEGO SUPERCOMPUTER CENTER

UC San Diego

```
install.packages('reshape2')  
library(reshape2)
```

long-to-wide: 'cast' repeated measure into wide table

```
W_cast = dcast(W_df,  
               formula = Date + Location + ... ~ WindGustDir,  
               fill = 0,  
               value.var = "WindGustSpeed")
```

WindGustDir distinguished the repeated measures

WindGustSpeed has the actual values

Extra to try:

```
# wide to long: ie 'melt' repeated measure into long  
# table
```

```
W_melt =melt(W_cast,  
             na.rm=TRUE,  
             measure.vars=c(23:39),  
             variable.name="WindGustDir_cast")
```

*The repeated measures are
now in columns 23 to 30*

pause

Reading Material

- **Data Preparation for Data Mining by Dorian Pyle**
 - http://www.ebook3000.com/Data-Preparation-for-Data-Mining_88909.html
- **Data mining – Practical Machine learning tools and techniques by Witten & Frank**
 - <http://books.google.com>

Many Variables

- **More variables \Rightarrow more information, but also more noise and more ways of interactions**
- **2 ways to handle many variables**
 - Variable Selection
 - Dimension reduction methods

Variable selection vs Dimensionality Reduction

- **Prior to algorithm, depends on data**
 - For large P , with noise particular to variables, try variable selection
 - For large P , diffuse noise, try dimension reduction by Matrix Factorization

Matrix Factorization:

Given a numeric matrix, can we reduce the number of columns?

conversely

Can we find interesting subspaces?

- **Matrix Factorization:**

Given a numeric matrix, can we reduce the number of columns?

- Yes, if features are constant or redundant

- **Matrix Factorization:**


Given a numeric matrix, can we reduce the number of columns?

- Yes, if features are constant or redundant
- Yes, if features only contribute noise
(conversely, want features that contribute to variations of the data)

Example: 2D data

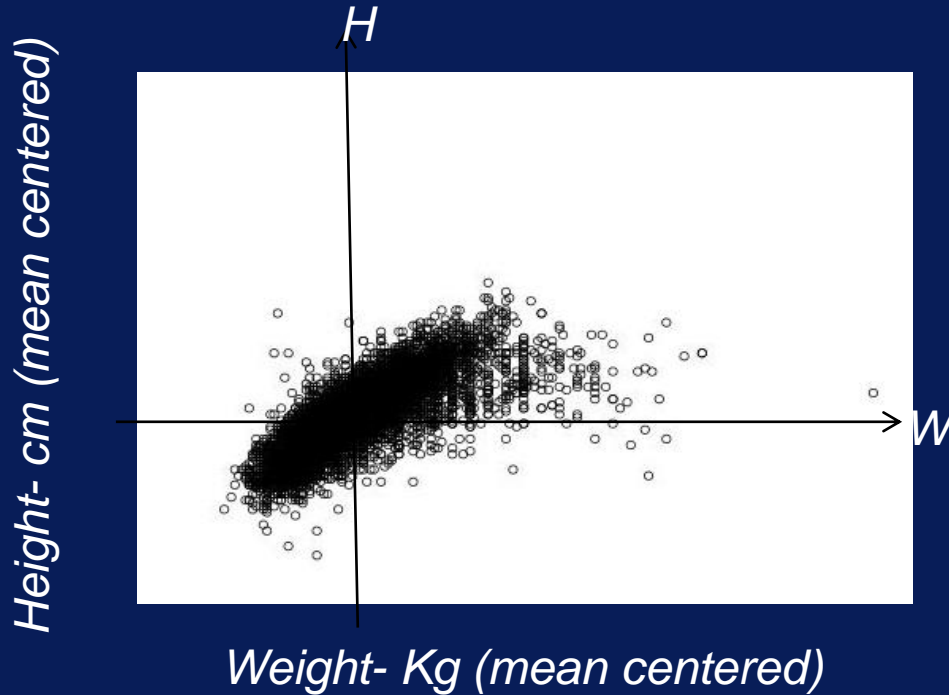
Weight Height

S1	50	179
S2	66	175
S3	74	180
S4	94	192




***this is
the
input
space***

Example: Athletes' Height by Weight

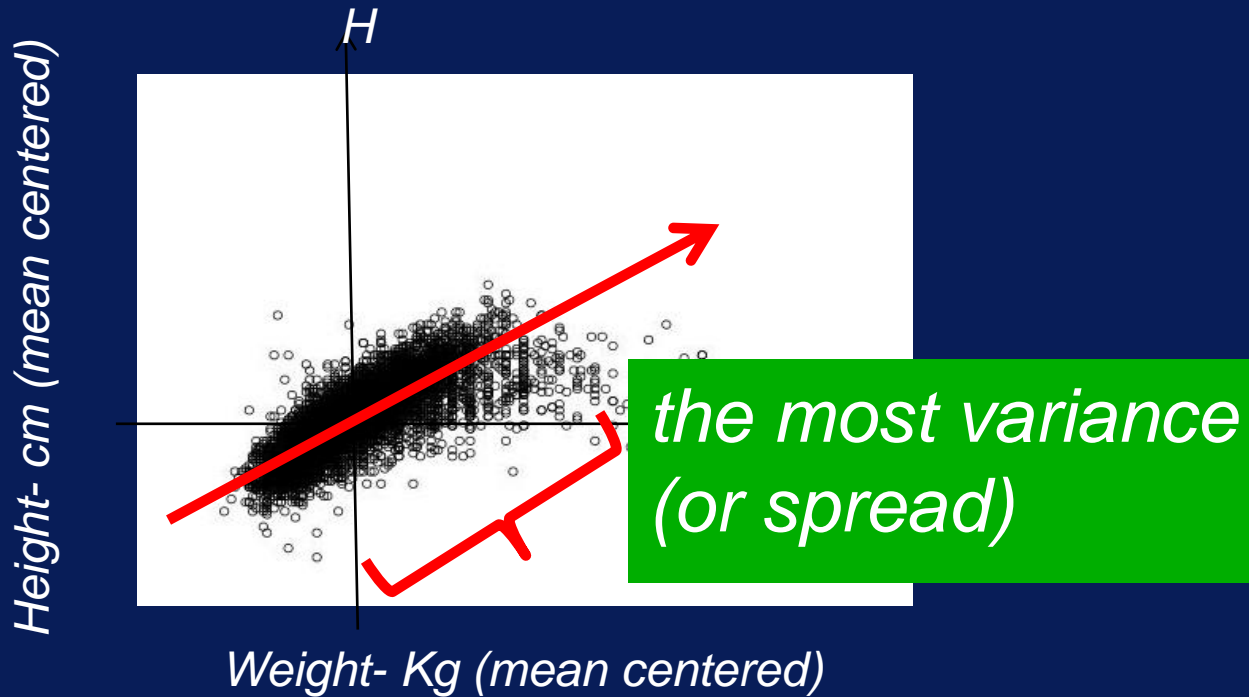


Find a line that aligns with the data.


Example: Athletes' Height by Weight




Find a line that aligns with the data.



Find a line that aligns with the data.




Find a line that aligns with the data.



Note that $W=1, H=0.8$ is a point on the line, for example.


Find a line that aligns with the data.



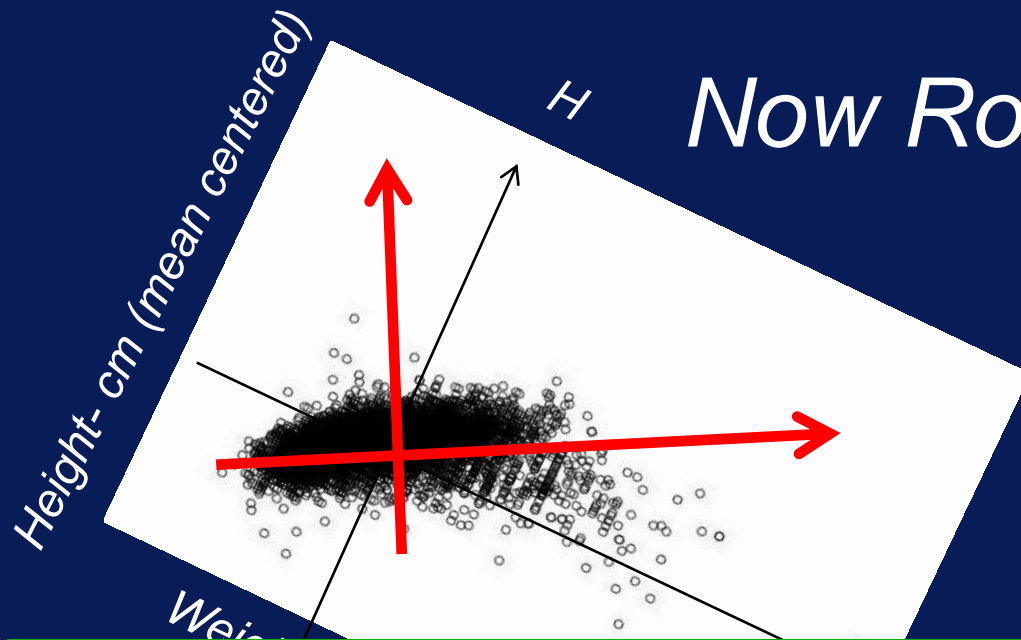
Note that $W=1, H=0.8$ is a point on the line, for example.

Let $[1 \ 0.8]$ represent the line, as a combination of W & H .

Find a line that aligns with the data.




The next direction of most variance.



Now Rotate Axis

*New axis (AKA features)
defined as combinations of
old features*

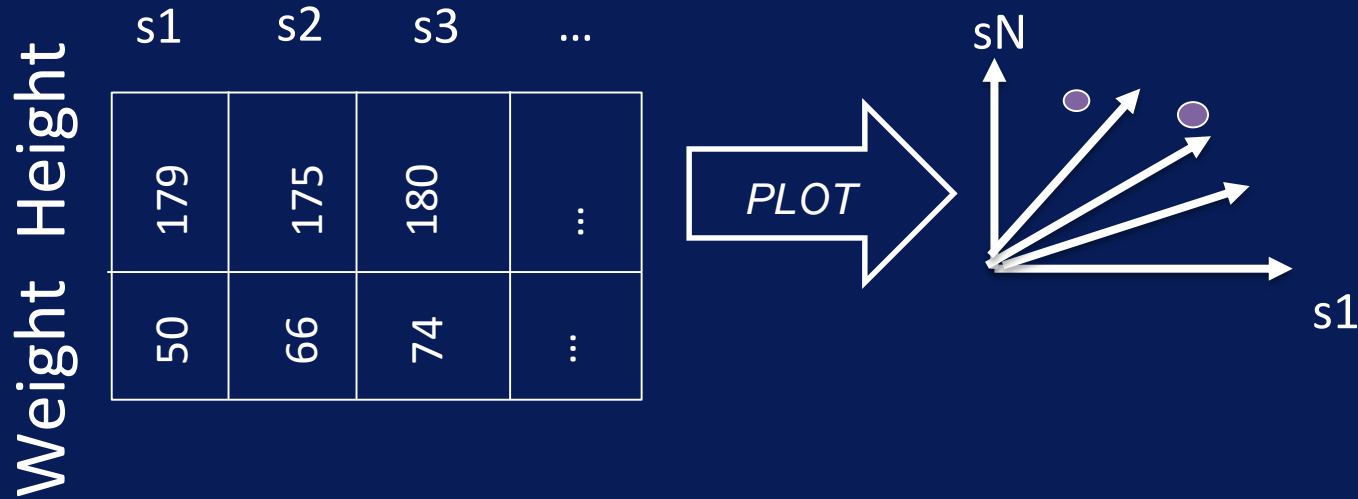
H *Possibly reduce dimensions*



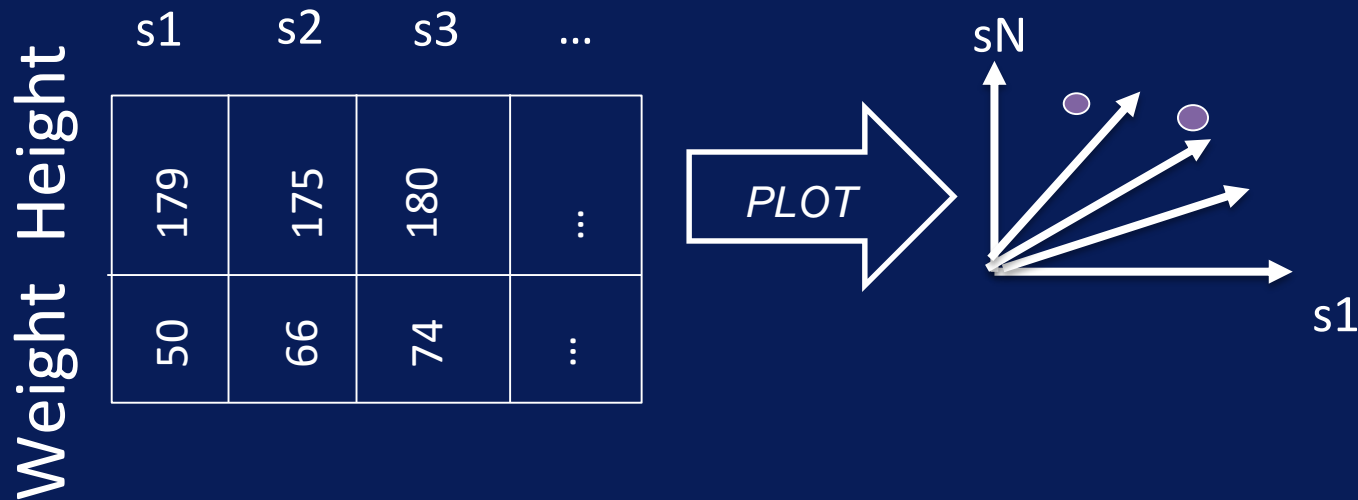
Project all points to one axis

(defined by the $[1 \ 0.8]$ 2D vector)

2D data transposed to 2 points in high dimensional space




2D data transposed to 2 points in high dimensional space



Same process as before, but now factors are N-dimensional vectors

Any Matrix can be approximated by outer product of factors



Original matrix




Factor A




Factor B


Any Matrix can be approximated by outer product of factors



Any Matrix can be approximated by outer product of factors



More factors gives better approximation




- Best Known Factorization Algorithms:
SVD (singular value decomposition)
PCA (principle component analysis)

- Best Known Factorization Algorithms:
SVD (singular value decomposition)
PCA (principle component analysis)

*Find orthogonal factors and
scale them down (i.e. normalize)*

SVD: factors and 'singular' scale values



- More generally:

Factorization Algorithms may vary depending on criterion for how factors 'align' with data.

- More generally:

Factorization Algorithms may vary depending on criterion for how factors 'align' with data.

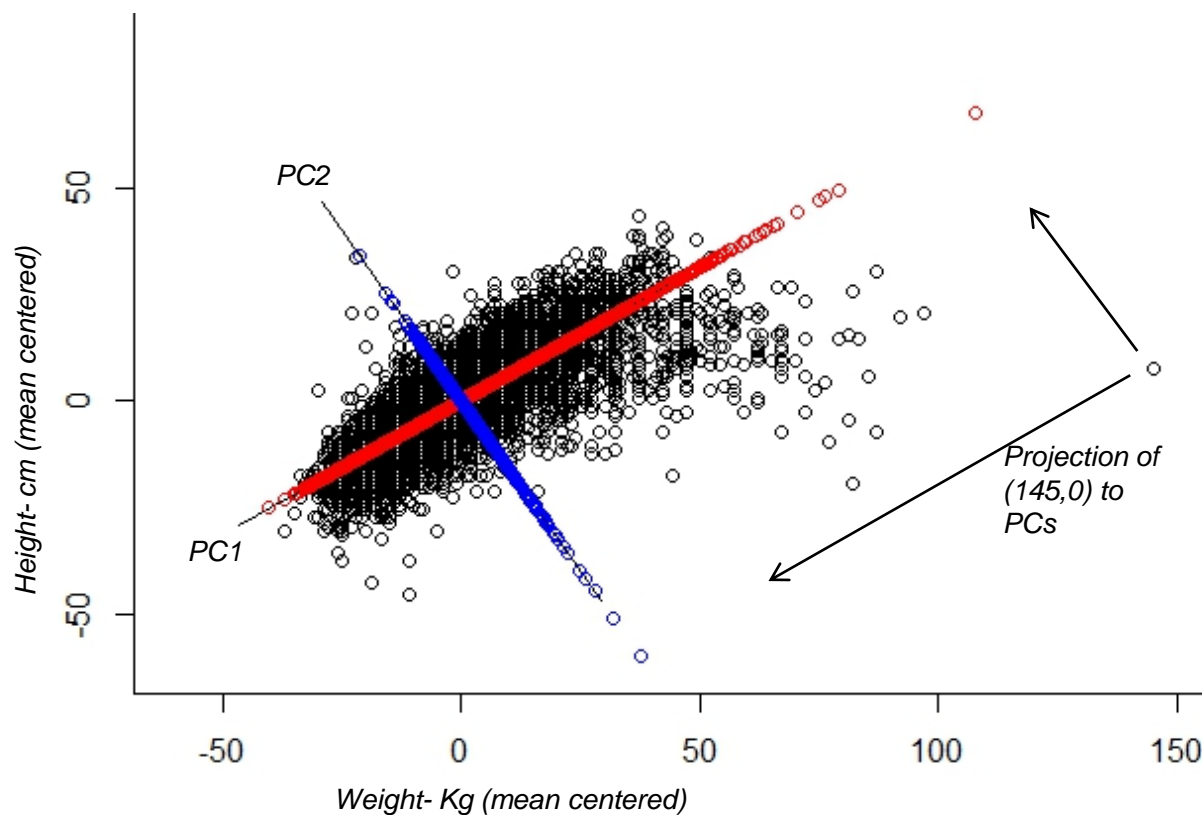
- Number of factors to use depends on tradeoff of good approximation vs good dimensional reduction

Can use cross validation or heuristics to choose.

Dimensionality Reduction via Principle Components

- **PCA:**
 - Find set of k vectors (aka factors) that describe data in alternative way
 - First component is the vector that maximizes the variance of data projected onto that vector
 - K -th component is orthogonal to all $k-1$ previous components

PCA conserves and reorders variance



Total Variance
Conserved:
 Var in Weight
+
 Var in Height
=
 $\text{Var in PC1} +$
 Var in PC2

In general:
 $\text{Var in PC1} >$
 $\text{Var in PC2} >$
 $\text{Var in PC3} \dots$

Principle Components

- Can choose k heuristically as approximation improves, or choose k so that 95% of data variance accounted
- aka Singular Value Decomposition
 - PCA on square matrices only
 - SVD gives same vectors on square matrices
- Works for numeric data only
- For higher dimensional data, use factors to visualize 2 factors at a time

Later, we'll compare SVD components with Clustering

#W_num is only numeric or integer fields of Weather data

```
> Wsvd=svd(W_num)
```

```
> str(Wsvd)
```

List of 3

```
$ d: num [1:9] 27442.7 231.2 96.4 68.2 44.5 ...
```

```
$ u: num [1:363, 1:9] -0.0524 -0.0521 -0.052 -0.0519 -0.0525 ...
```

```
$ v: num [1:9, 1:9] -0.005042 -0.014276 -0.000969 -0.00314 -0.005491 ...
```

- end