Other ML related activites

Complex Social Science Gateway – a tool for cross-cultural analysis in R

Select dataset,
Select varialbes,
Submit analysis

http://socscicompute.ss.uci.edu/

(but moving soon)

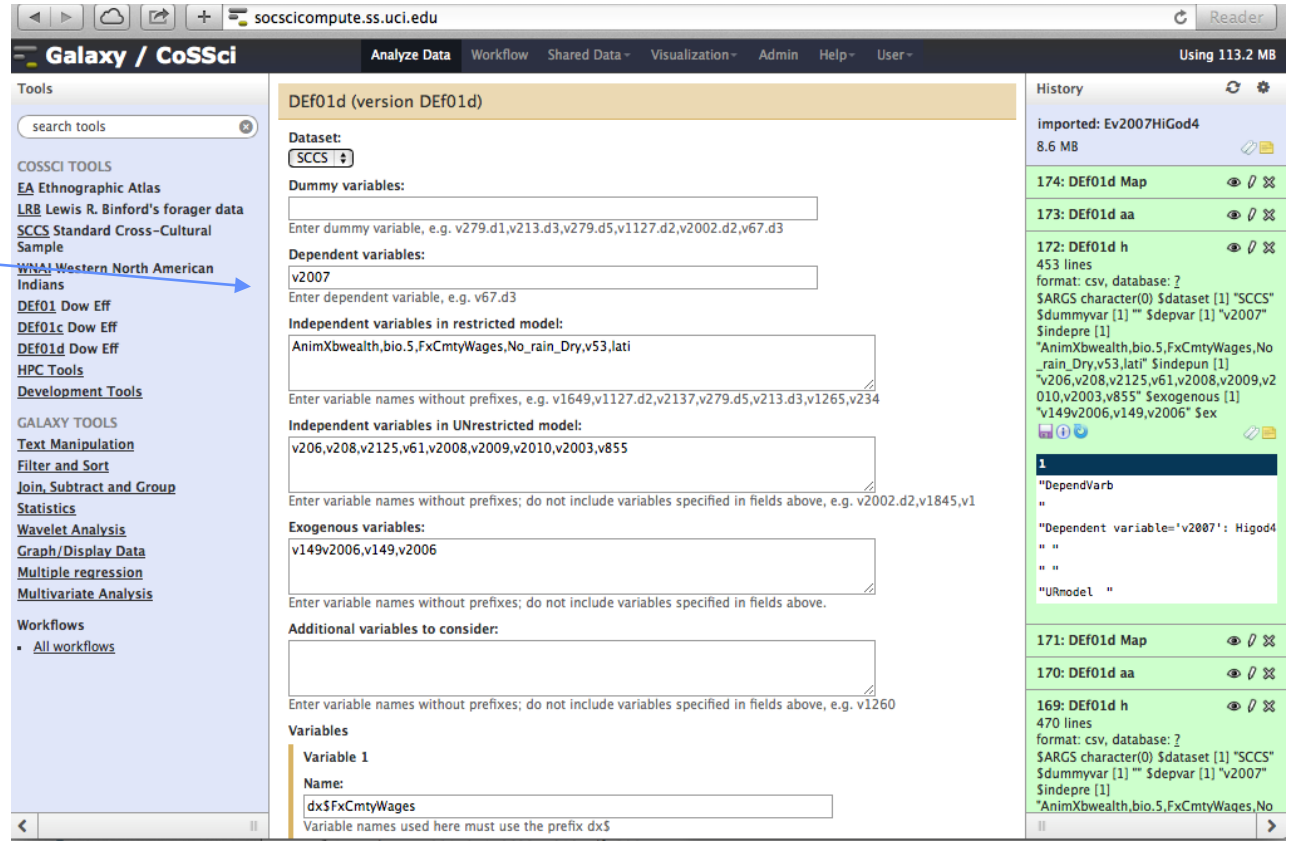

R Analysis options

- **Two-stage least squares to handle spatially correlated errors (OLS, logit, multinomial logit)**

- **Bootstrap sampling of Bayesian network (package bnlearn) to confirm OLS effects, or suggest other moderating/mediating effects**

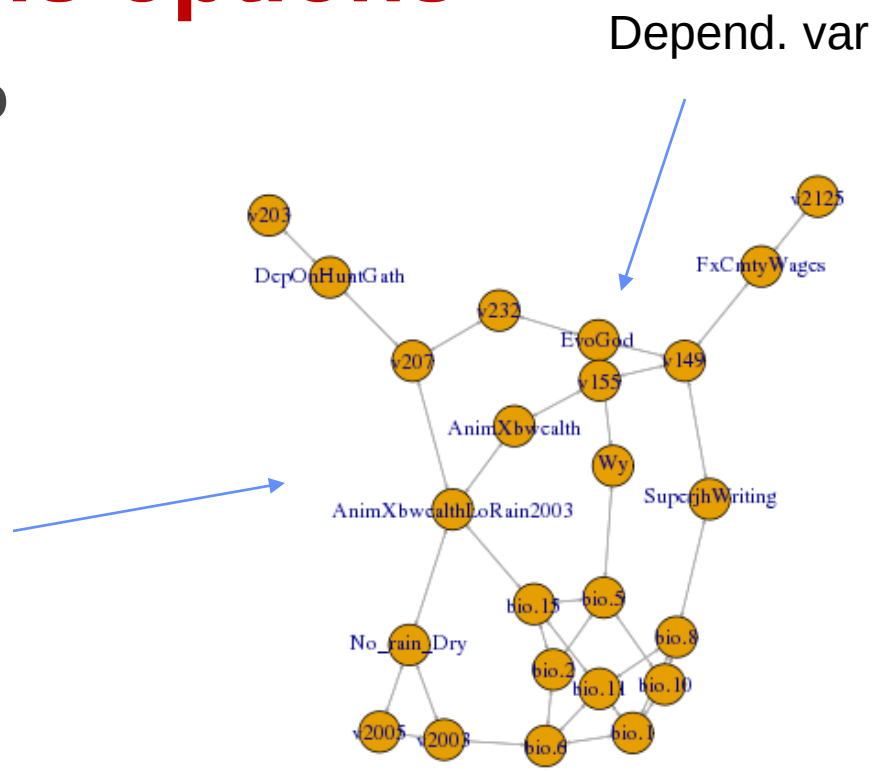

In a nutshell: Bayesian Network captures probabilistic dependencies between variables

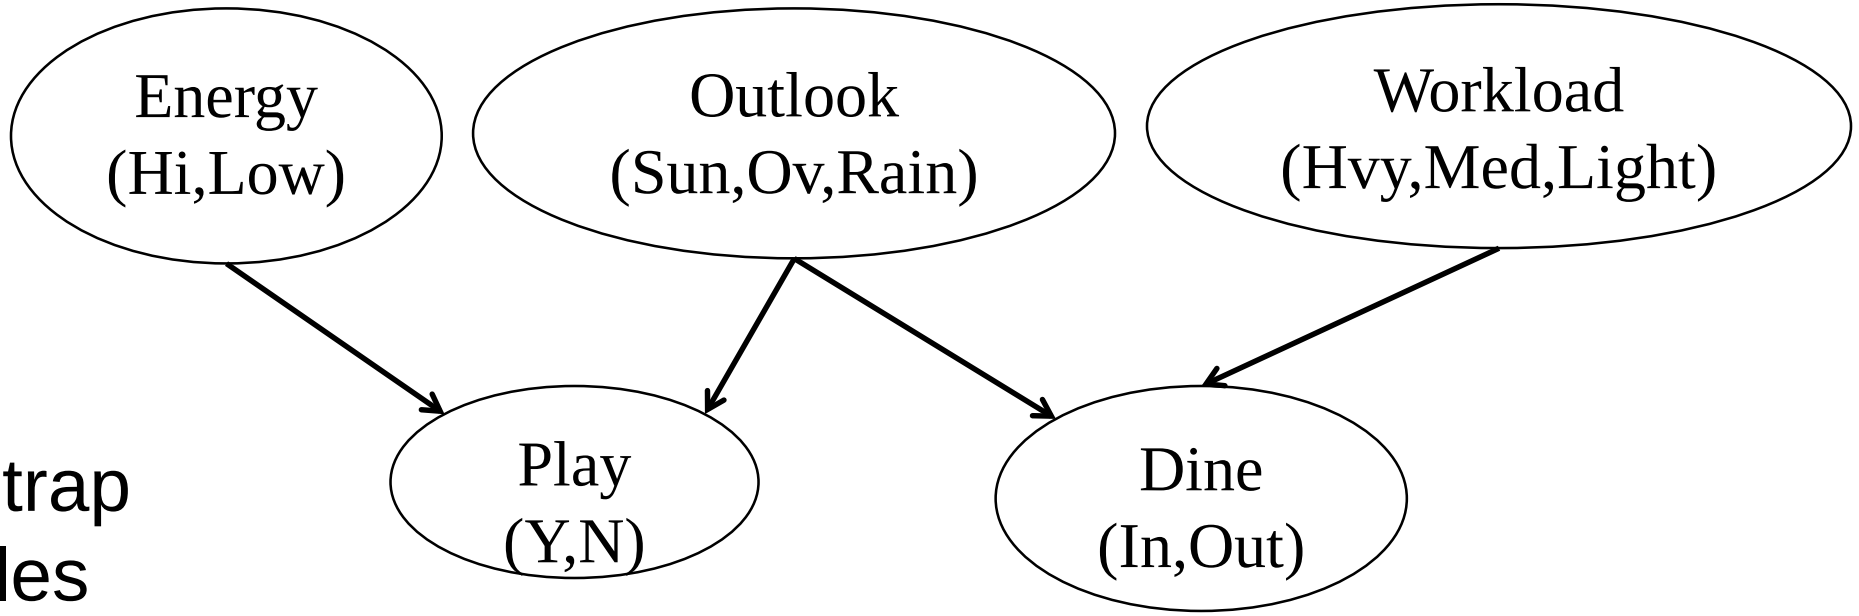

AND:
Bootstrap samples determine most likely edges

Image Analysis of Rural Photography 175K war and depression era photos extracting features for datamining

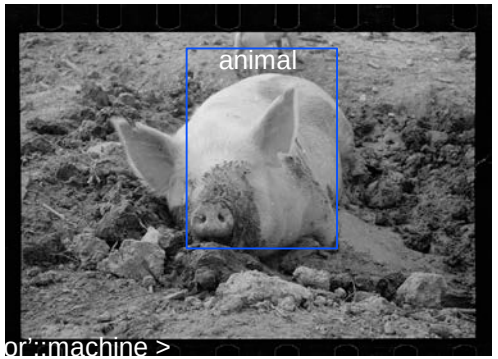

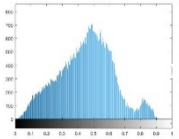

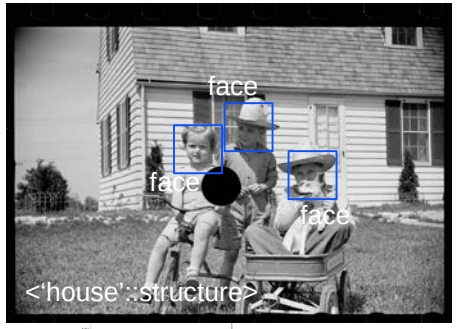

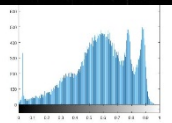

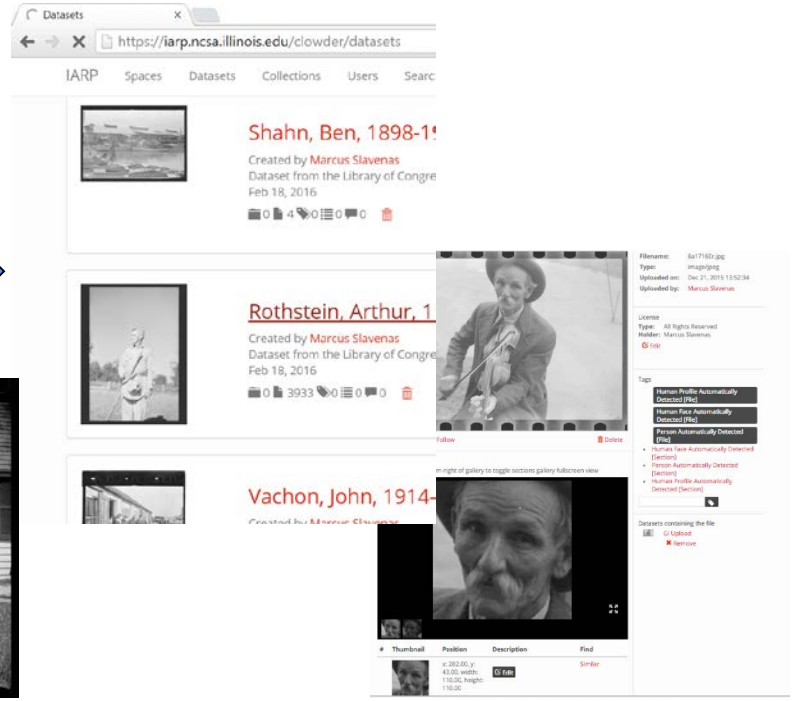

Title:

”Destitute pea pickers in California. Mother of seven children.”

Histogram of Gradients

CellSize = [32 32]

Feature length = 10260

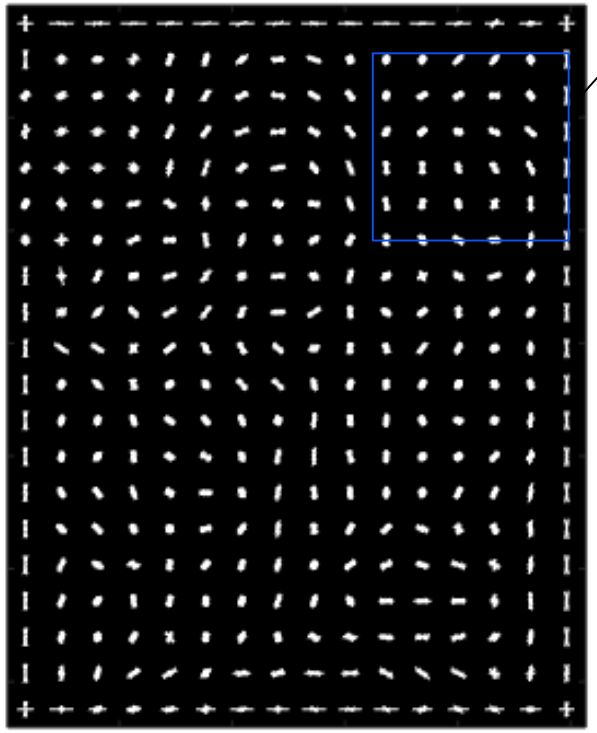

For each pixel in a cell, take filters:

[-1 0 1]

And

[-1

0

1]

Take weighted average and bin into 9 orientations; the bin frequency is like magnitude

Title:

”Destitute pea pickers in California.
Mother of seven children.”

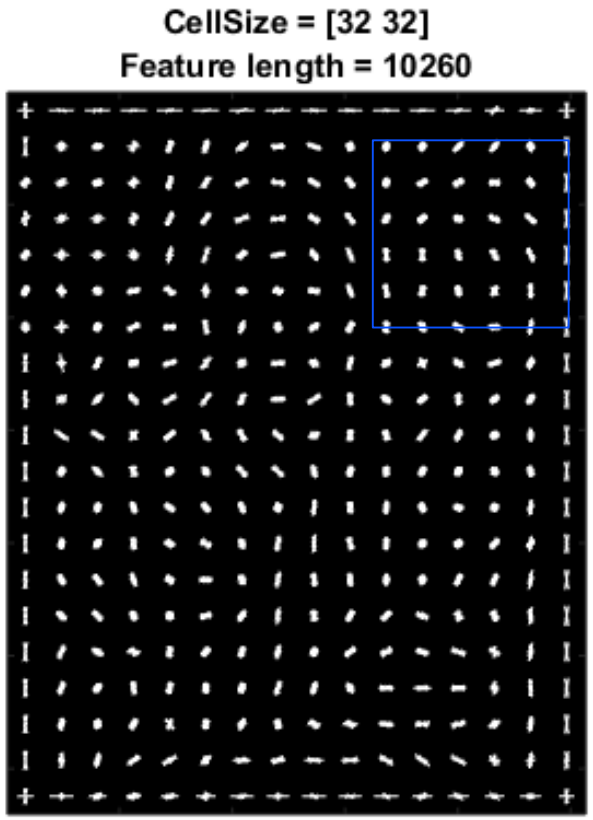

Take all orientations, at different scales, as 1 big vector, and feed into classifier trained to recognize Face.

Title:

”Destitute pea pickers in California. Mother of seven children.” By D. Lange, 1936, California, [metadata]

Metadata processing:

- **Parse and tag speech (using Stanford NLP tools, word ontologies, in Python NLP toolkit)**
- **Several words identify ‘person’**

- **SQL: give me all pictures by Lange with possible ‘person’ and num_faces > 0**

Early 20th century, ~15k prison Bertillon id cards extracting information

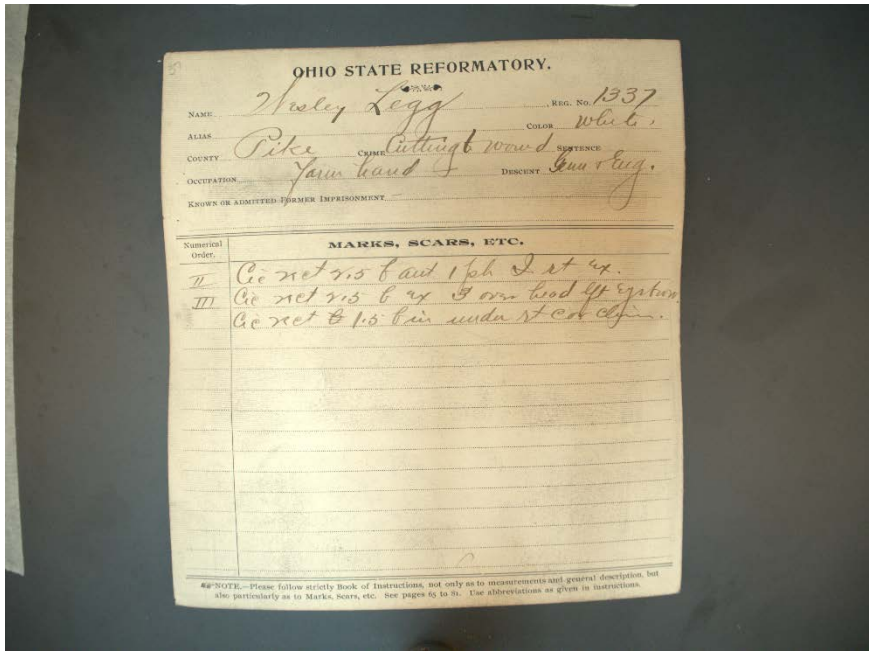

Segment,
binarize,
denoise

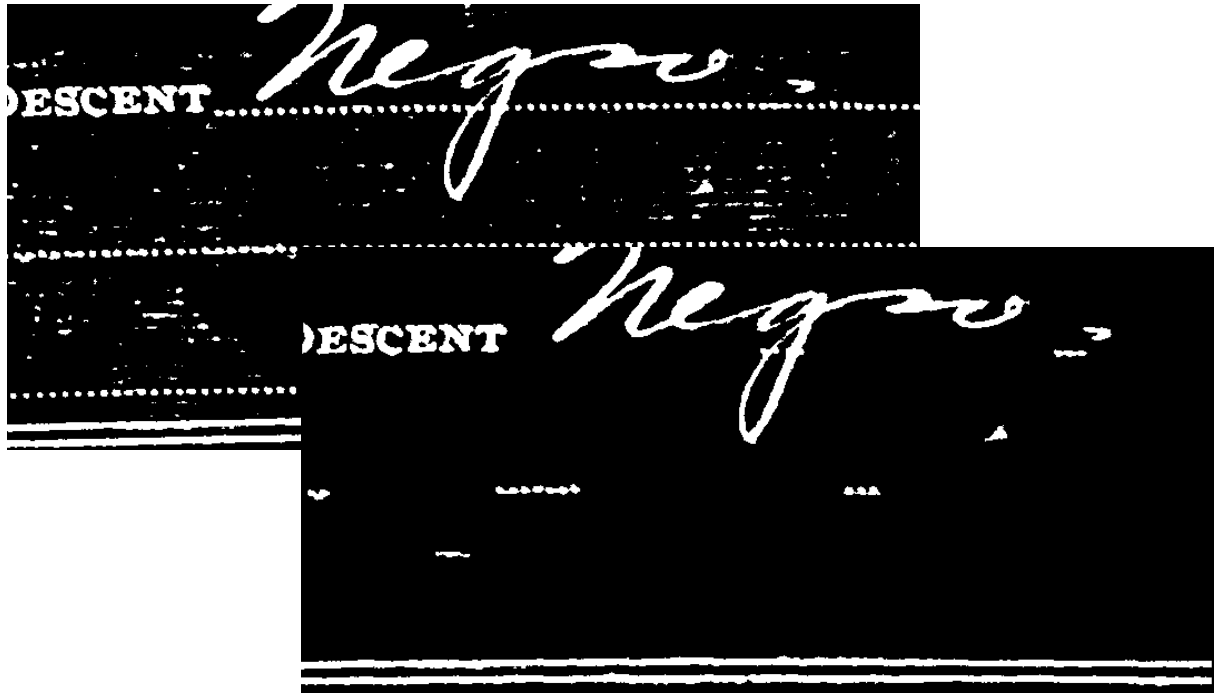

extract field and cell

(word spotting) Get profile and
compare to known templates

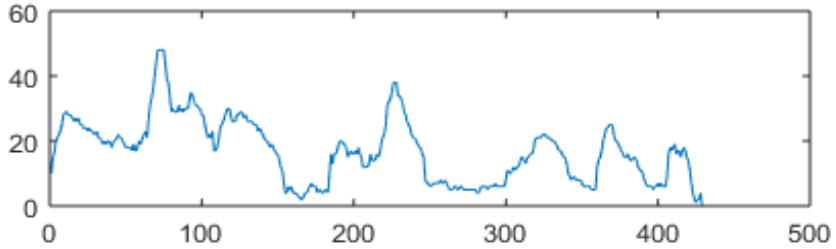

Linear to Logistic to Neural Network model

- *yi*= $bo * 1 + b_1 * xi_1$ + $b_2 * xi_2$ … = B*X

- **Squash** $bo * 1 + b_1 * xi_1$ **to 0,1 range using Logistic Function:**

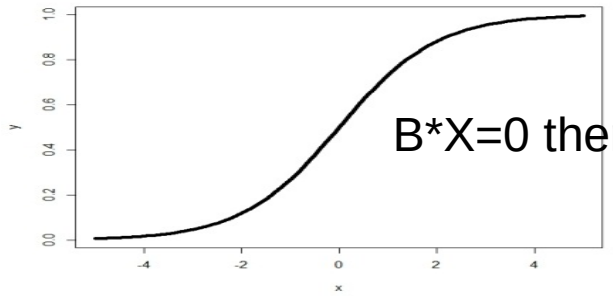

Logistic Regression to Neural Networks

- **Use several squash functions (hidden layer)**

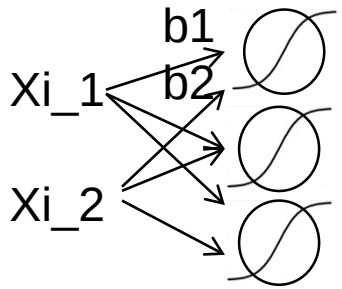

- **Take further combinations (output layer)**

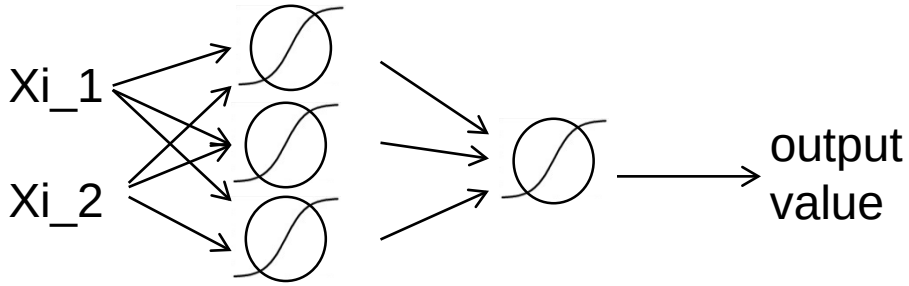

- **More powerful but more complex**

  many parameters, many options, needs more training

organize connections into cells, add layers (deepen), add special pooling operations at some layers - you get a convolution network

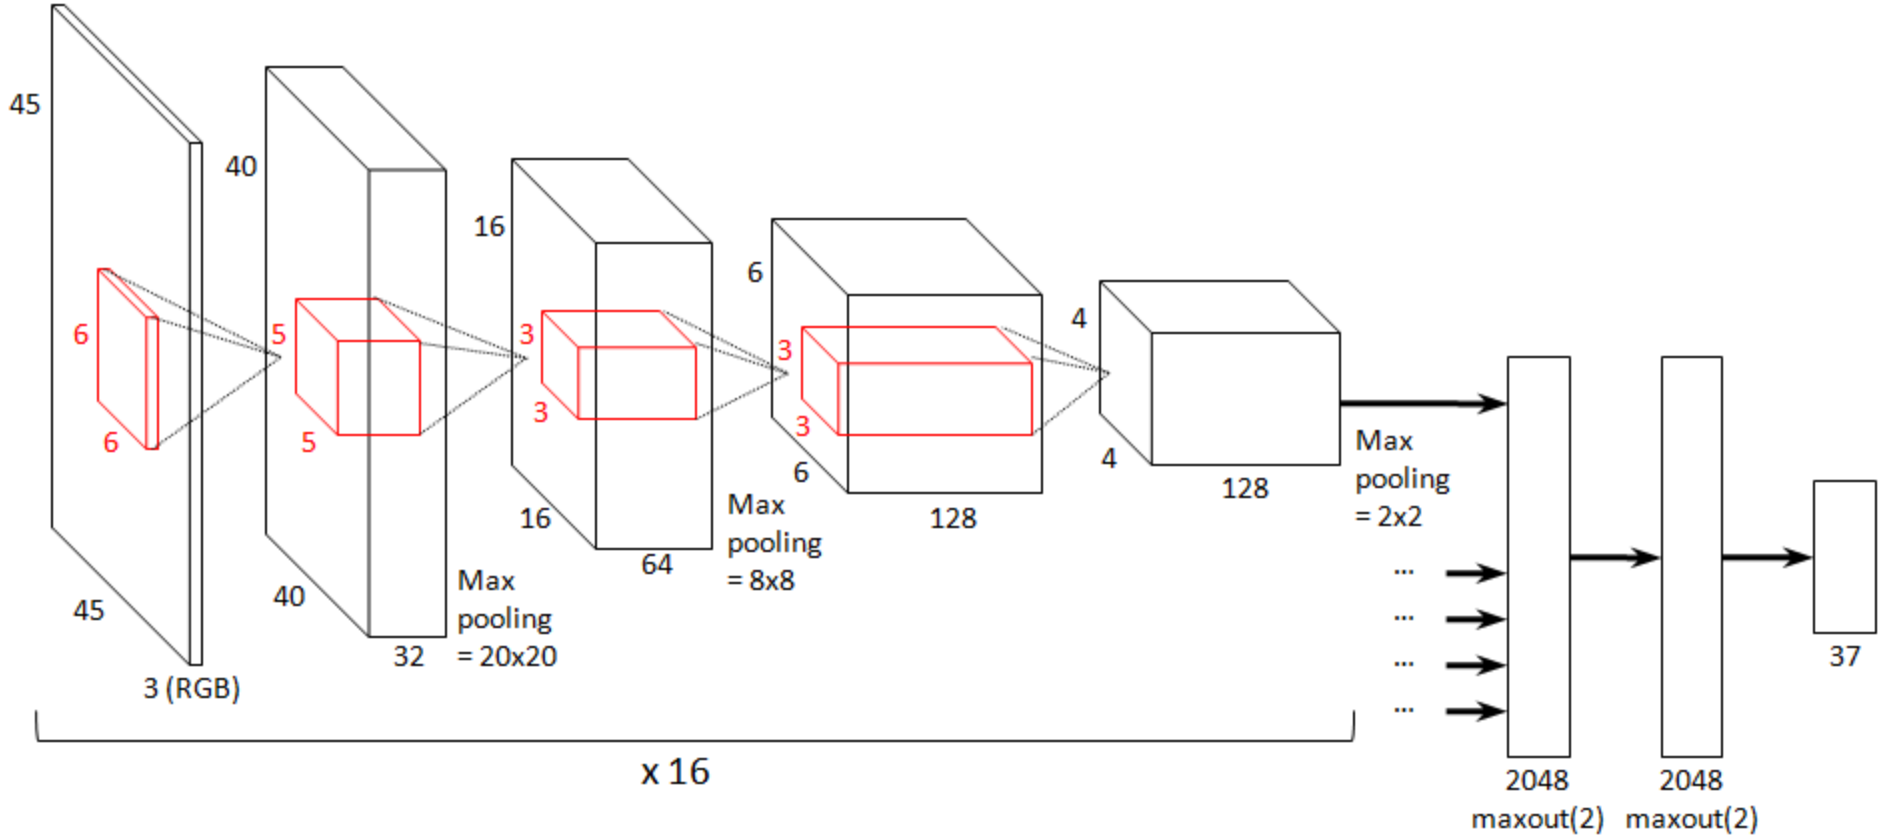

SciKit python package has a convolution neural network

```
nn2 = Classifier(
        layers=[
         Convolution("Rectifier", channels=numch, kernel_shape=(10,10),pool_shape=(2,2)),
         Convolution("Rectifier", channels=numch, kernel_shape=(6,6),pool_shape=(4,4)),
         Layer("Sigmoid",units=numalpha2do*4),
         Layer("Sigmoid",units=numalpha2do*2)
         ],
         verbose=False,
    learning_rate=0.001,valid_set=(Xtrain,Ytrain),
    n_iter=myiter)
    nn2.fit(Xtrain,Ytrain)
```

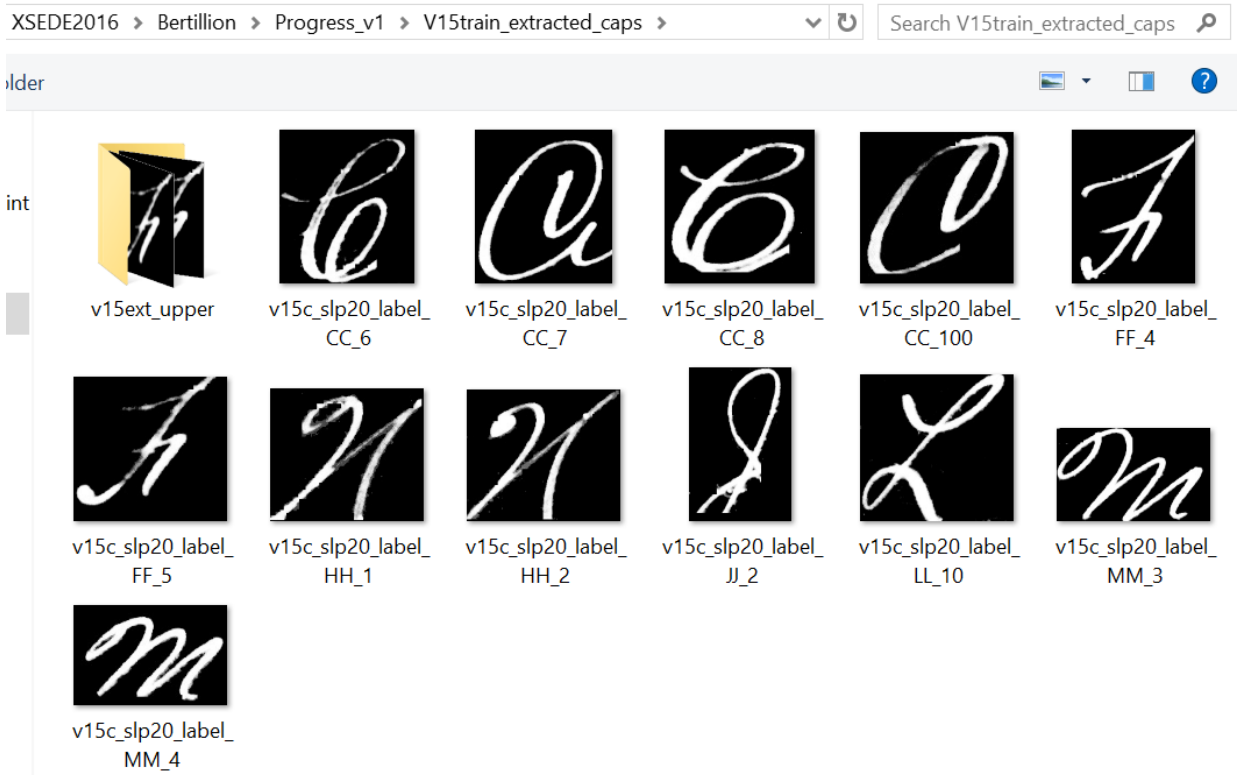
XSEDE2016 › Bertillion › Progress_v1 › V15train_extracted_caps ›
Search V15train_extracted_caps
v15ext_upper
v15c_slp20_label_CC_6
v15c_slp20_label_CC_7
v15c_slp20_label_CC_8
v15c_slp20_label_CC_100
v15c_slp20_label_FF_4
v15c_slp20_label_FF_5
v15c_slp20_label_HH_1
v15c_slp20_label_HH_2
v15c_slp20_label_JJ_2
v15c_slp20_label_LL_10
v15c_slp20_label_MM_3
v15c_slp20_label_MM_4

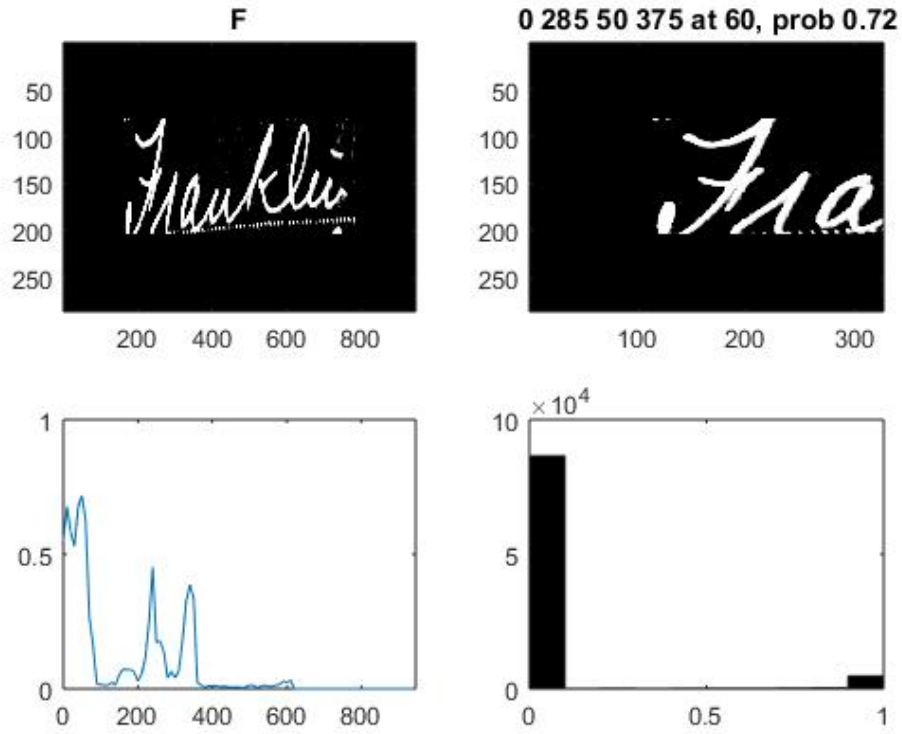
F
0 285 50 375 at 60, prob 0.72
50
100
150
200
250
200 400 600 800
100 200 300
1
0.5
0
0 200 400 600 800
×10^4
10
5
0
0 0.5 1

Topic Modelling with Latent Dirichlet Allocation

- Each circle is 1 word occurrence
- 2 topics (filled/empty circles), 15 docun
- Initially random assignments

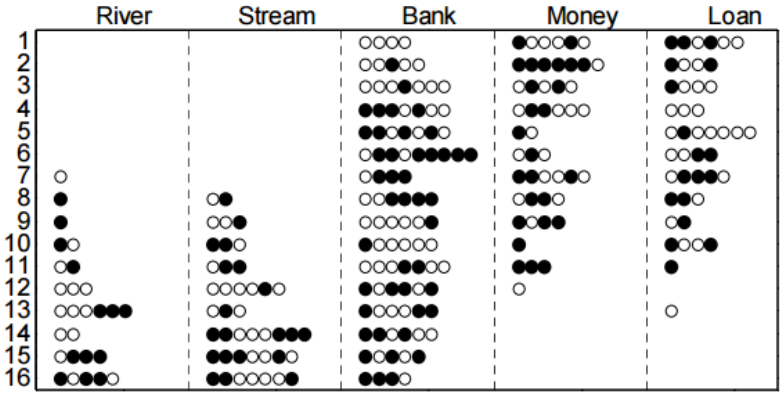

- After learning, topics are well formed

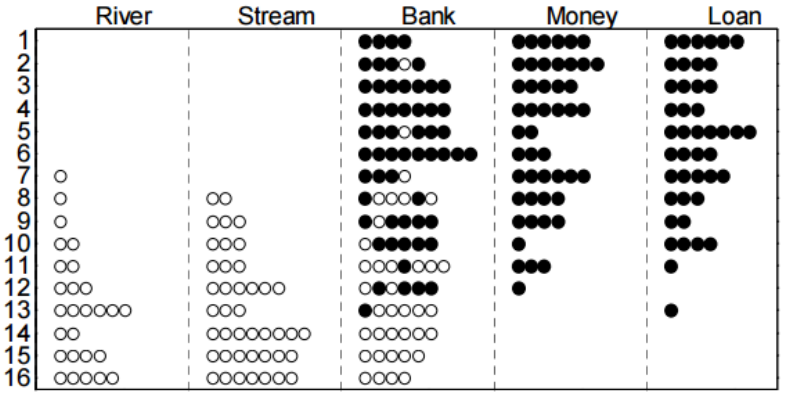

LDA optimization

- **Start with initial guess of topic=*t*, and parameters**
- **Repeat:**

  Compute the expected value of word=*w*

  Compute the parameters that maximize likelihood L of *t given w*

Parameters are estimated from word/topic counts

With each iteration, objective function L goes up

Topic Modelling with Latent Dirichlet Allocation on HPC

- R LDA package: wraps C programs for Gibbs sampling or EM
- Mallet: Gibbs sampling multicore, java code
- Spark LDA: EM
- Asymptotic Distributed LDA : MPI based, no bells&whistles

- **Example Case:**

articles from post WWII journals

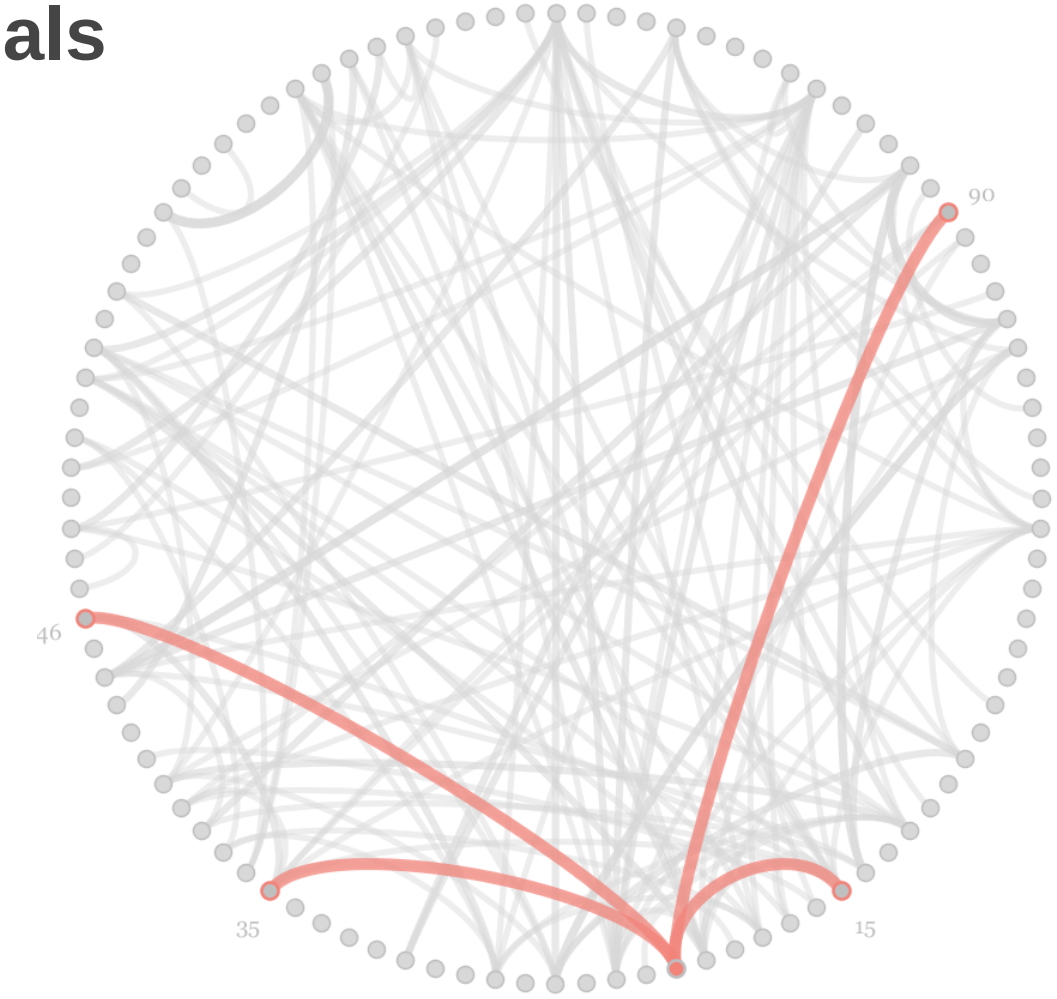

topic21	topic35	topic46	topic90	topic15
worker	county	farm	house	work
labor	city	land	family	day
employment	state	farmer	area	labor
job	town	acre	home	time
percent	york	crop	city	pay
industry	public	agricultural	build	hour
defense	district	county	live	week
service	mayor	family	room	wage
work	relief	agriculture	unit	make
increase	local	cotton	community	month
employ	population	labor	project	condition
train	person	state	income	year
unemployment	citizen	area	rend	employ
wage	community	migrant	move	find
employee	place	rural	neighborhood	case
department	residence	year	facility	money
employer	resident	tenant	urban	service
occupation	board	migration	low	receive
rate	settlement	make	occupy	leave
number	large	large	resident	care
industrial	welfare	grower	lodge	employer
production	number	small	condition	order
earnings	part	california	apartment	good
woman	aid	camp	neighborhoods	require

- **Example Case:**

Sample topic plot (tree map)