

Scaling R options



R on Comet strengths

- **Some examples what R is especially good for that you might want to do on Comet:**
- **Sampling/bootstrap methods,**
- **Data Wrangling,**
- **Particular Statistical procedures that you won't find implemented anywhere else, e.g.**
 - Multiple Imputation methods,
 - Instrument Variable (2 stage) Regression
 - Matching observational data subjects

R on Comet

At Unix prompt:

- *module load R*
- *R* (for interactive R session)

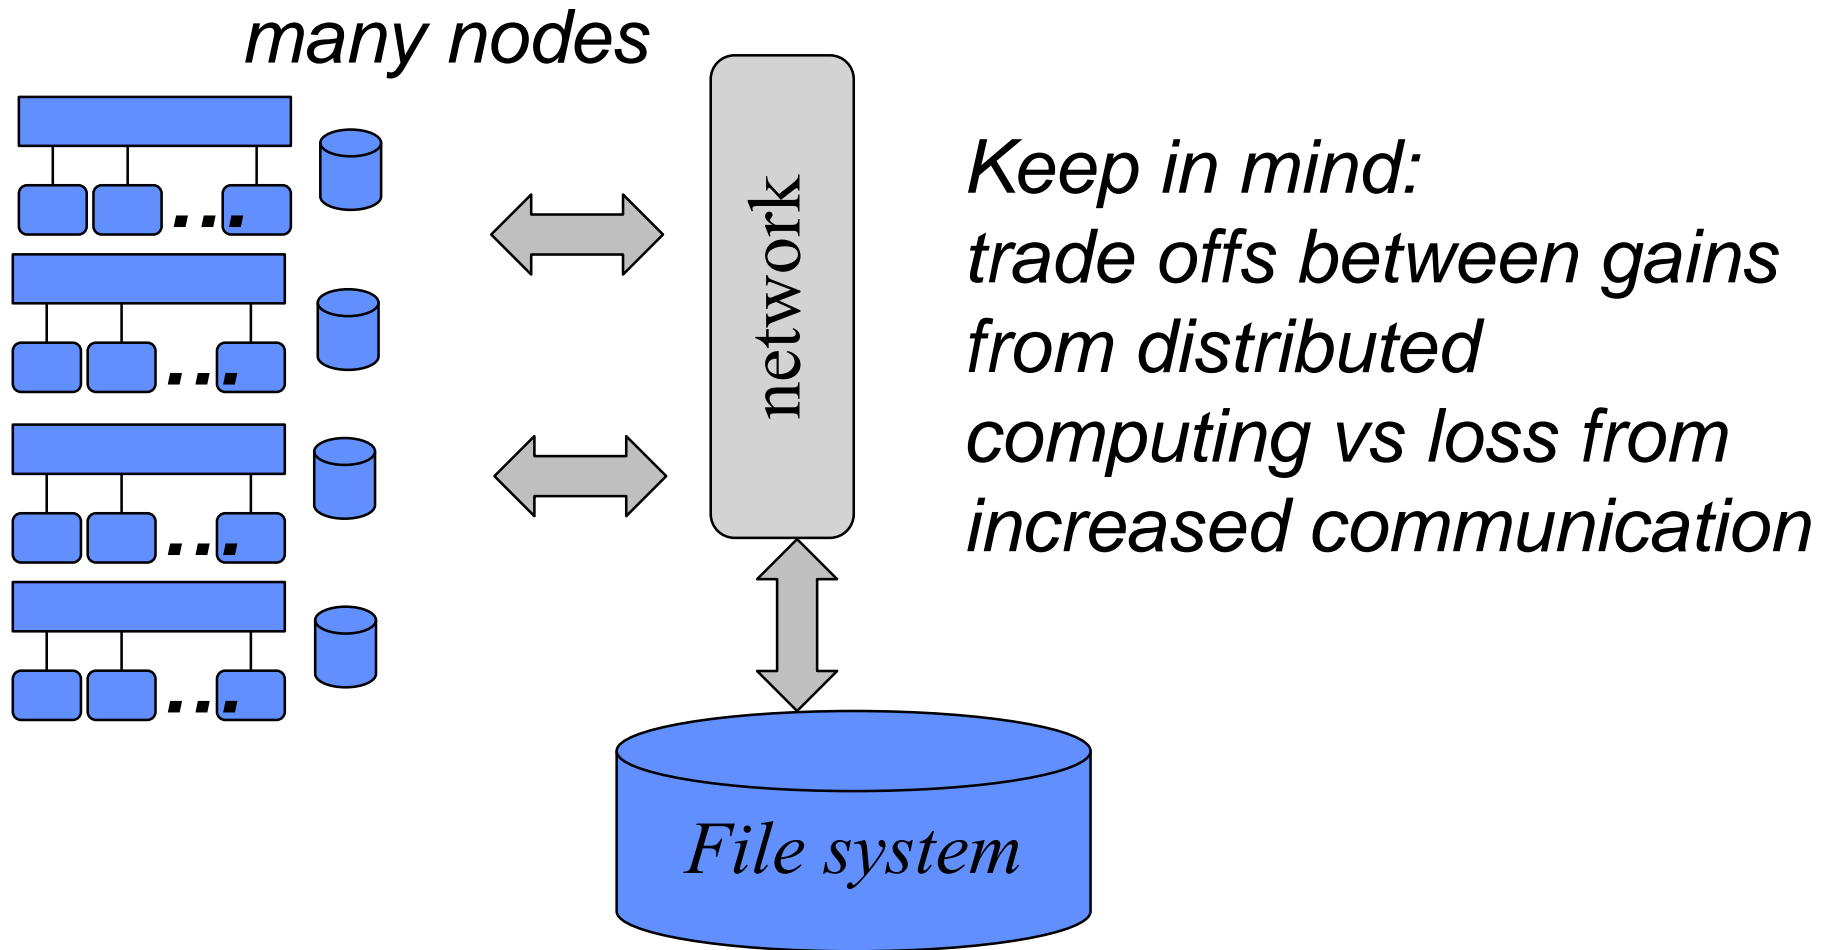
or

- *Rscript script-name* (for batch)

Scaling, practically

- **Scaling (with or without more data):**
 - more complex analysis (ie optimizations)
 - more sampling (ie more trees in Random Forest)
- **Sometimes easy to parallelize (like with sampling),**
- **Sometimes too much communication between parts (matrix inversion)**

On Distributed Computing

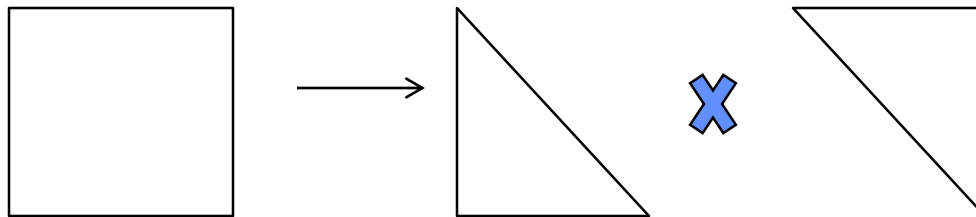


Scaling In a nutshell

- **R takes advantage of math libraries for vector operations**
- **R packages provide multicore, multinode (snow), or map/reduce (RHadoop) options**
- **However, model implementations not necessarily built to use parallel backends**
 - Some models more amenable to parallel versions

Consider Regression Computations

- **Linear Model:** $Y = X * B$
where Y =outcomes , X =data matrix
- **Algebraically, we could:**
 - take “inverse” of $X * Y = B$ (time consuming)
 - use derivatives to search for solutions (very general)
- **Or, better:**
 - QR decomposition of X into triangular matrices (easier to solve but more memory)



Consider Regression models in R

- **Related Models and Functions :**

lm() #Linear Model

glm() #Generalized Linear Model
(logistic regression, etc)

aov() #Analysis of Variance
(returns ANOVA table of F-scores)

All these work on system of equations

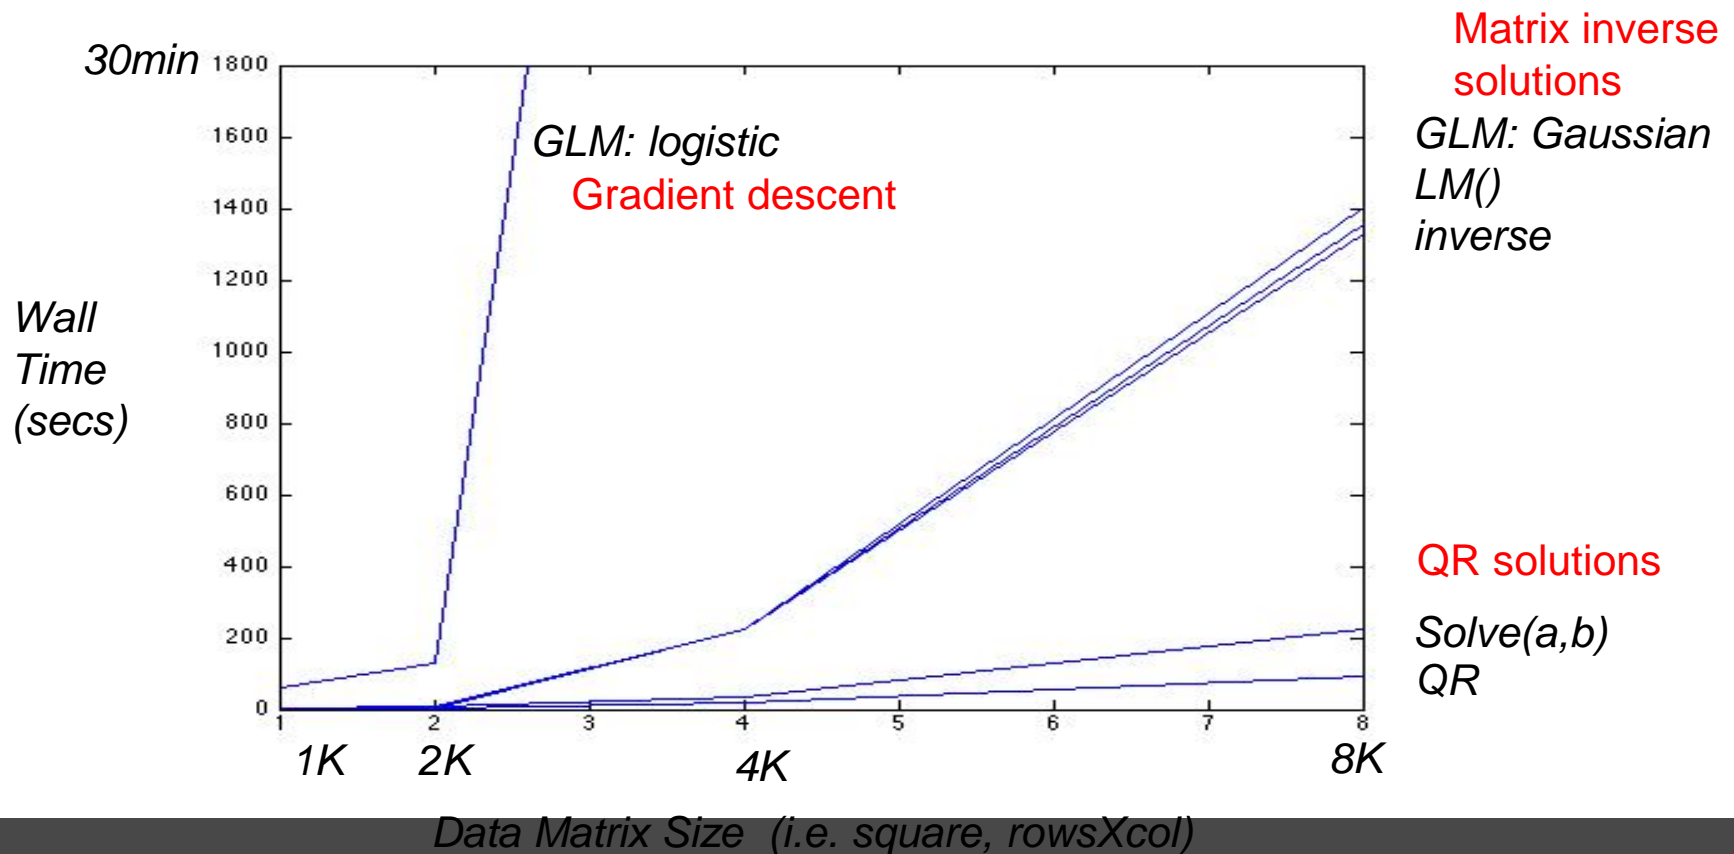
Solving Linear Systems

Performance with R, 1 compute node

R:

`glm(Y~X,family=gaussian)` #gaussn regrssn (like lm)

`glm(Y~X,family=binomial)` # logistic regrssn (Y=0 or 1)



R multicore

- **Intel Math Kernel Libraries provides fast operations for vector operations**
- **Uses threads across cpu cores to pass data & commands**

R multicore

- Run loop iterations on separate cores

```
install.packages(doMC)  
library(doMC)  
registerDoMC(cores=24)  
getDoParWorkers()
```

allocate workers

*%dopar% puts loops
across cores,
(loops are independent)
%do% runs it serially*

```
results = foreach(i=1:24,.combine=rbind) %dopar%  
{ ... your code here
```

return(a variable or object)

*specify to combine results into
array with row bind*

*returned items
'combined' into list
by default*

R multinode: parallel backend

- Run loop iterations on separate nodes

```
install.packages('doSNOW')
library('doSNOW')
...
cl <- makeCluster( mpi.universe.size()-1, type='MPI' )
clusterExport(cl,c('data'))
registerDoSNOW(cl)

results = foreach(i=1:47,.combine=rbind) %dopar%
{ ... your code here

      return( a variable or object )
})
stopCluster(cl)
mpi.exit()
```

*allocate cluster as
parallel backend*

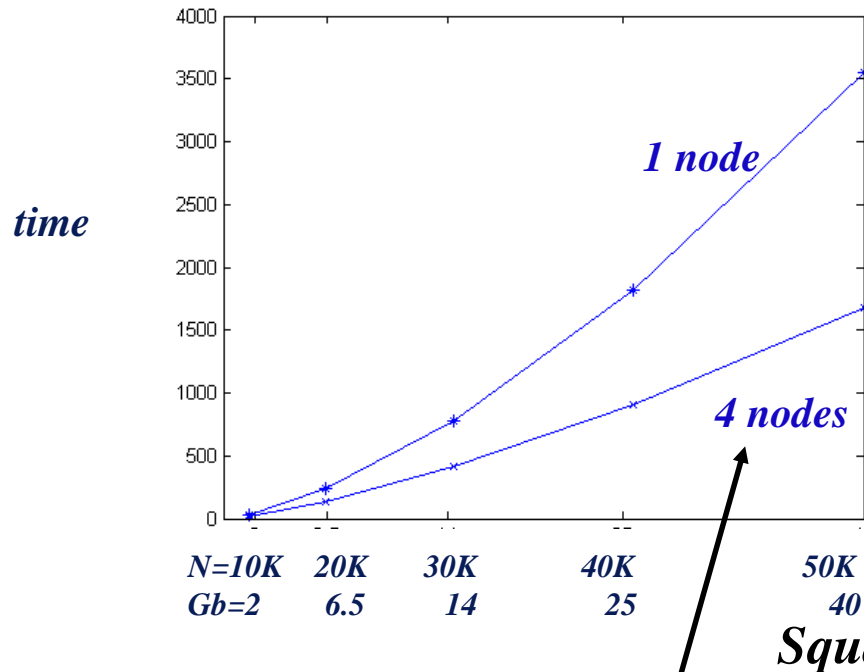
*%dopar% puts loops
across cores and
nodes*



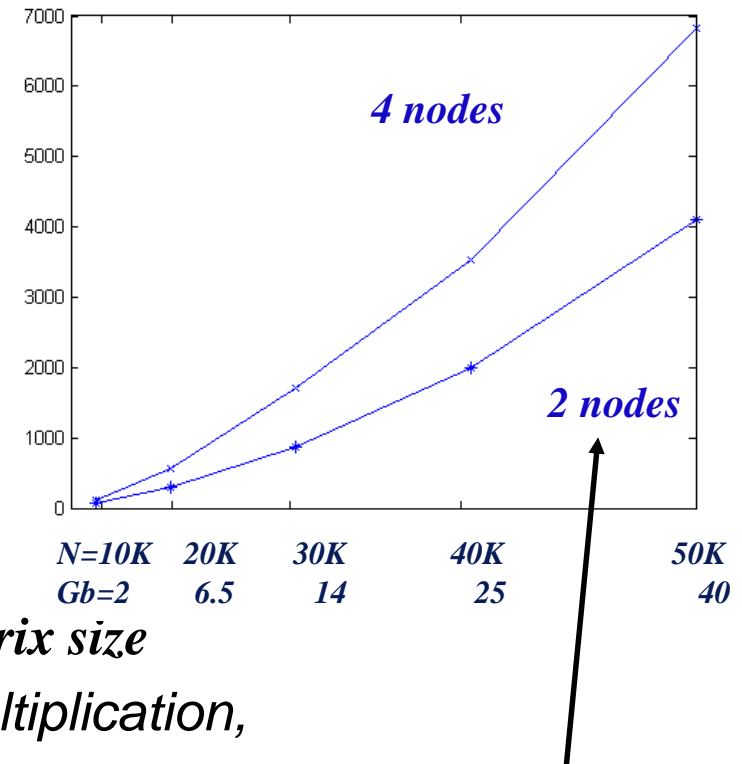
Multiple Compute Nodes not always help

(tested on Gordon)

Matrix Multiplication



Matrix Inversion



multinodes: more nodes is less time for multiplication,

less nodes is better for inversion

Another Parallel option:

- **Serially packing R jobs onto cores**
 1. **batch job and calls MPI**
 2. **MPI executes a Perl script on each core**
 3. **Perl script gets cpu-id and passes it to R**
 4. **R uses cpu-id to process some particular input**

Normal
batch
job info

```
train100@comet-ln3:~/Rtrain/Rpacking_serial
[train100@comet-ln3 Rpacking_serial]$ more comet_sbatch_serial_packed
#!/bin/bash
# -----
# slurm script for a batch job on comet
# to run a task on individual cores
# -----
#SBATCH --job-name="serial-pack"
#SBATCH --output="serial-pack.%j.%N.out"
#SBATCH --partition=shared
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=4
#SBATCH --export=ALL
#SBATCH -t 00:30:00

bash

#Generate a hostfile from the slurm node list
export SLURM_NODEFILE=`generate_pbs_nodefile`

#Run job from working directory or do something like this:
#cd /oasis/scratch/comet/$USER/mydirectory/

module load R

#mpirun executes the bundler perl script on each core, the bundler will launch
# R with arguments to indicate which file to process
# NOTE the argument to bundler is the number of R tasks to execute
#       it should be = ntask-per-node X nodes    (but >= would work)
mpirun_rsh -hostfile $SLURM_NODEFILE -np 24 ./bundler.pl 48
[train100@comet-ln3 Rpacking_serial]$
```

mpirun the
'bundler' perl
script on 24
nodes

the
'bundler'
script

The argument
'was 48 tasks
to do

Get current
cpu id and
number of
processes

Depending on
the cpu-id,
execute R
and pass an
input file

```
train100@comet-06-55:Rpacking_serial
#!/usr/bin/perl
use strict;
use warnings;
# -----
# Perl 'wrapper' script to help launch tasks (ie a program)
# This script is called by:
#   mpirun_rsh -np N ./bundler.pl num_tasks
#   which will put an instance of this script on each core
# -----
my $ntasks          = $ARGV[0];
my ($myid, $numprocs) = split(/\s+/, `./getid`); #getid will return cpu id and total processor in MPIjob
# -----
# using getid, launch only a task in turn
#   pass use the task id as an argument
#   (note there can be more tasks than processors)
# -----
for (my $i=0; $i<$ntasks; $i++) {
    if($myid == $i % $numprocs) {
        print "Perl Bundler: $myid -th Task\n";
        my $filei_2use=$i+1;
        System("/opt/R/bin/Rscript R_LinModel_randXY.R $filei_2use");
    }
}
~
~
```

22, 2-9 All

Packing Serial with large Random Forest job

- Option 1: Run separate trees on separate cores

```
install.packages(doMC)
registerDoMC(cores=15)
getDoParWorkers()
library("randomForest");
results = foreach(i=1:15,.combine=rbind) %dopar%
  {RF1 <-randomForest(formula,data=X,na.action=na.omit,
    importance=TRUE,
    ntree=100000,
    do.trace=1,
    nodesize=1)
  classRF1$importance
}
```

allocate workers ←

%dopar% puts loops across cores, (loops are independent) %do% runs it serially ↙

returned items 'combined' into list ↗

Sampling on large data could be huge ↘

Packing Serial with large Random Forest job

- **Option 2: split tree sampling to make it embarrassingly parallel**

ie run R script on separate cores and average results

Can speed up processing without losing interesting variable combinations

R and Map/Reduce

- Map/Reduce streaming can pipe Input/Output through R
- R can make direct map/reduce calls with R-Hadoop interface (from Revolution Analytics)

Hadoop Map/Reduce Interfaces with R

(slides from G.Lockwood SDSC)

- R Streaming (simplest) or Hadoop API
- E.g. streaming Hadoop is like piping input/output

```
cat input | Rscript mapper.R | sort | Rscript reducer.R > output
```



You provide these two scripts; Hadoop does the rest

Paradigmatic Example: Word Counting

How would you count all the words in Moby Dick?

Call me Ishmael. Some years ago - never mind how long precisely - having little or no money in my purse, and nothing particular to interest me on shore, I thought I would sail about a little and see the watery part of the world. It is a way I have of driving off the spleen and regulating the circulation.

How could you count all the words in all web pages? (assume the data is spread out over many nodes)

Use Map/Reduce, take computation to nodes

Wordcount: Hadoop streaming mapper

Emit key-value pairs ('cat' is 'concatenate and print')

```
emit.keyval <- function(key, value) {  
  cat(key, '\t', value, '\n', sep='')  
}
```

```
stdin <- file('stdin', open='r')  
while ( length(line <- readLines(stdin, n=1)) > 0 ) {
```

Split line
Into words

```
  line <- gsub('(^\\s+|\\s+$)', '', line)  
  keys <- unlist(strsplit(line, split='\\s+'))  
  value <- 1
```

Use words
as keys

```
  lapply(keys, FUN=emit.keyval, value=value)  
}  
close(stdin)
```

Example from Glen Lockwood, SDSC

What One Mapper Does

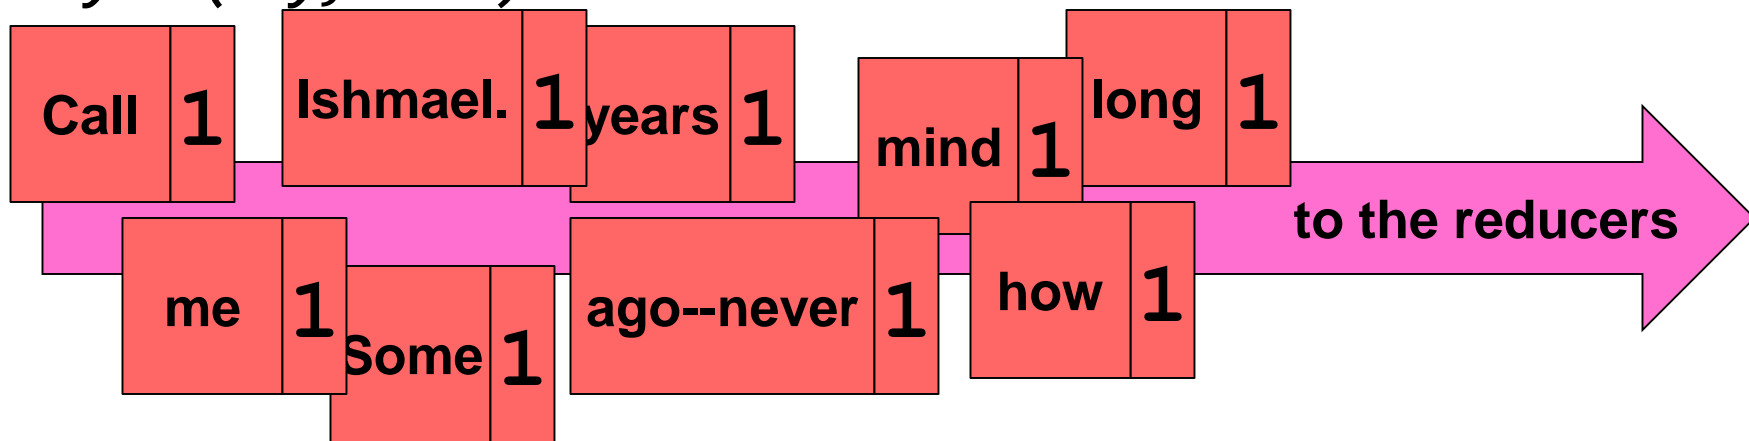
line =

Call me Ishmael. Some years ago—never mind how long

keys =

Call	me	Ishmael.	Some	years	ago--never	mind	how	long
------	----	----------	------	-------	------------	------	-----	------

emit.keyval(key, value) ...



Reducer Loop

- **If this key is the same as the previous key,**
 - add this key's value to our running total.
- **Otherwise,**
 - print out the previous key's name and the running total,
 - reset our running total to 0,
 - add this key's value to the running total, and
 - "this key" is now considered the "previous key"

Wordcount: Streaming Reducer (1/2)

```
last_key <- ""
running_total <- 0

stdin <- file('stdin', open='r')
while ( length(line <- readLines(stdin,n=1)) > 0 ) {
  line <- gsub('^\\s+|\\s+$',' ', line)
  keyvalue <- unlist(strsplit(line, split='\\t', fixed=TRUE))
  this_key <- keyvalue[[1]]
  value <- as.numeric(keyvalue[[2]])

  if ( last_key == this_key ) {
    running_total <- running_total + value
  }
  else {
    (to be continued...)
  }
}
```

Get key,
Value

Add up
values

Wordcount: Streaming Reducer (2/2)

For each new
key, emit
<key, sum>

```
else {  
    if ( last_key != "" ) {  
        cat(  
paste(last_key, '\t', running_total, '\n', sep='') )  
        }  
        running_total <- value  
        last_key <- this_key  
    }  
}  
  
if ( last_key == this_key ) {  
    cat( paste(last_key, '\t', running_total, '\n', sep='') )  
}  
  
close(stdin)
```

Testing Mappers/Reducers

- Debugging Hadoop is not fun

```
$ head -n100 pg2701.txt |  
  ./wordcount-streaming-mapper.R | sort |  
  ./wordcount-streaming-reducer.R  
  
...  
with 5  
word, 1  
world. 1  
www.gutenberg.org 1  
you 3  
You 1
```

Launching Hadoop Streaming

```
$ hadoop dfs -copyFromLocal ./pg2701.txt mobydick.txt
```

```
$ hadoop jar \  
/opt/hadoop/contrib/streaming/hadoop-streaming-1.0.3.jar \  
-D mapred.reduce.tasks=2 \  
-mapper "Rscript $PWD/wordcount-streaming-mapper.R" \  
-reducer "Rscript $PWD/wordcount-streaming-reducer.R" \  
-input mobydick.txt \  
-output output
```

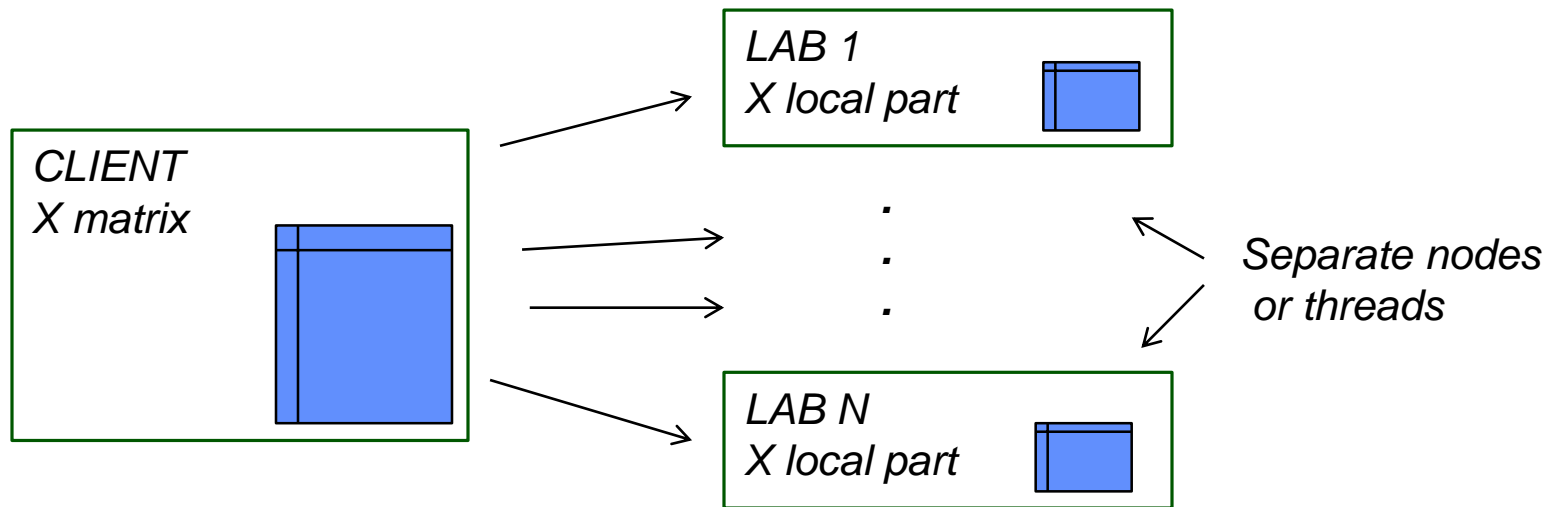
```
$ hadoop dfs -cat output/part-* > ./output.txt
```

Ask XSEDE support for latest Hadoop scripts

Matlab quickview

- **Distributed Toolbox:**

- allocate distributed matrices using 'spmd' code
- MPI or threads under the hood
- You decide data/task set up



SPMD Matlab – general scheme

%Matlab code using SPMD toolbox, where X is a matrix

...

spmd

*Distribute Matrix X
and process local part*

fprintf('in 1st spmd %i \n',labindex);

Xsd=codistributed(X, codistributor('1d', 1)); %distrib rows

XsdLoc=getLocalPart(Xsd);

end

% Other code...

spmd

local-result ~ calculations on XsdLoc

*It stays distributed
in memory for further
processing within SPMD block*

end;

gather local-result

% ... rest of code....

- **Some future stuff**

Other R packages:

- **Rspark** - R interface to Spark
- **pdbR** - higher level over R-MPI, distributed matrix support and other
- **HiPLAR** - GPU and multicore for linear algebra
- **Rgputools** – GPU support
- **R openMP** , better data mgt than dopar, parallel (mclapply)
- **Ff**, **bigmemory**; **Revolution Scale R** (commercial) – map data to files

- **R install.packages on Comet**
- **Need to specify url and yes to personal library;**
- **If compiling is required you might get an error**

Login into Comet,

cd to Rtrain/Rmulticore

module load R

R

```
[screen 0: bash] etrain102@comet-03-09:~/SI2016ML/Rmulticore
[etrain102@comet-03-09 Rmulticore]$
[etrain102@comet-03-09 Rmulticore]$
[etrain102@comet-03-09 Rmulticore]$ module load R
[etrain102@comet-03-09 Rmulticore]$ R

[p4rodrig@comet-ln3 Rtrain]$ rm -r etrain102/
[p4rodrig@comet-ln3 Rtrain]$
```

Exercise: R multicore

Login into Comet,

cd to Rtrain/Rmulticore

#Step 1 generate random data

#Step 2, setup sampling #Step 3 set up parallel environment

#Step 3 - -----Parallel Options -----

```
library('doMC');
```

```
registerDoMC(23)    # < ---- number of cores here
```

Now run a bootstrap sample of regression

```
coef_boot=vector('numeric',num_resamp) #an empty vector to save results
```

Use 'do' or 'dopar' to test single/multicore execution

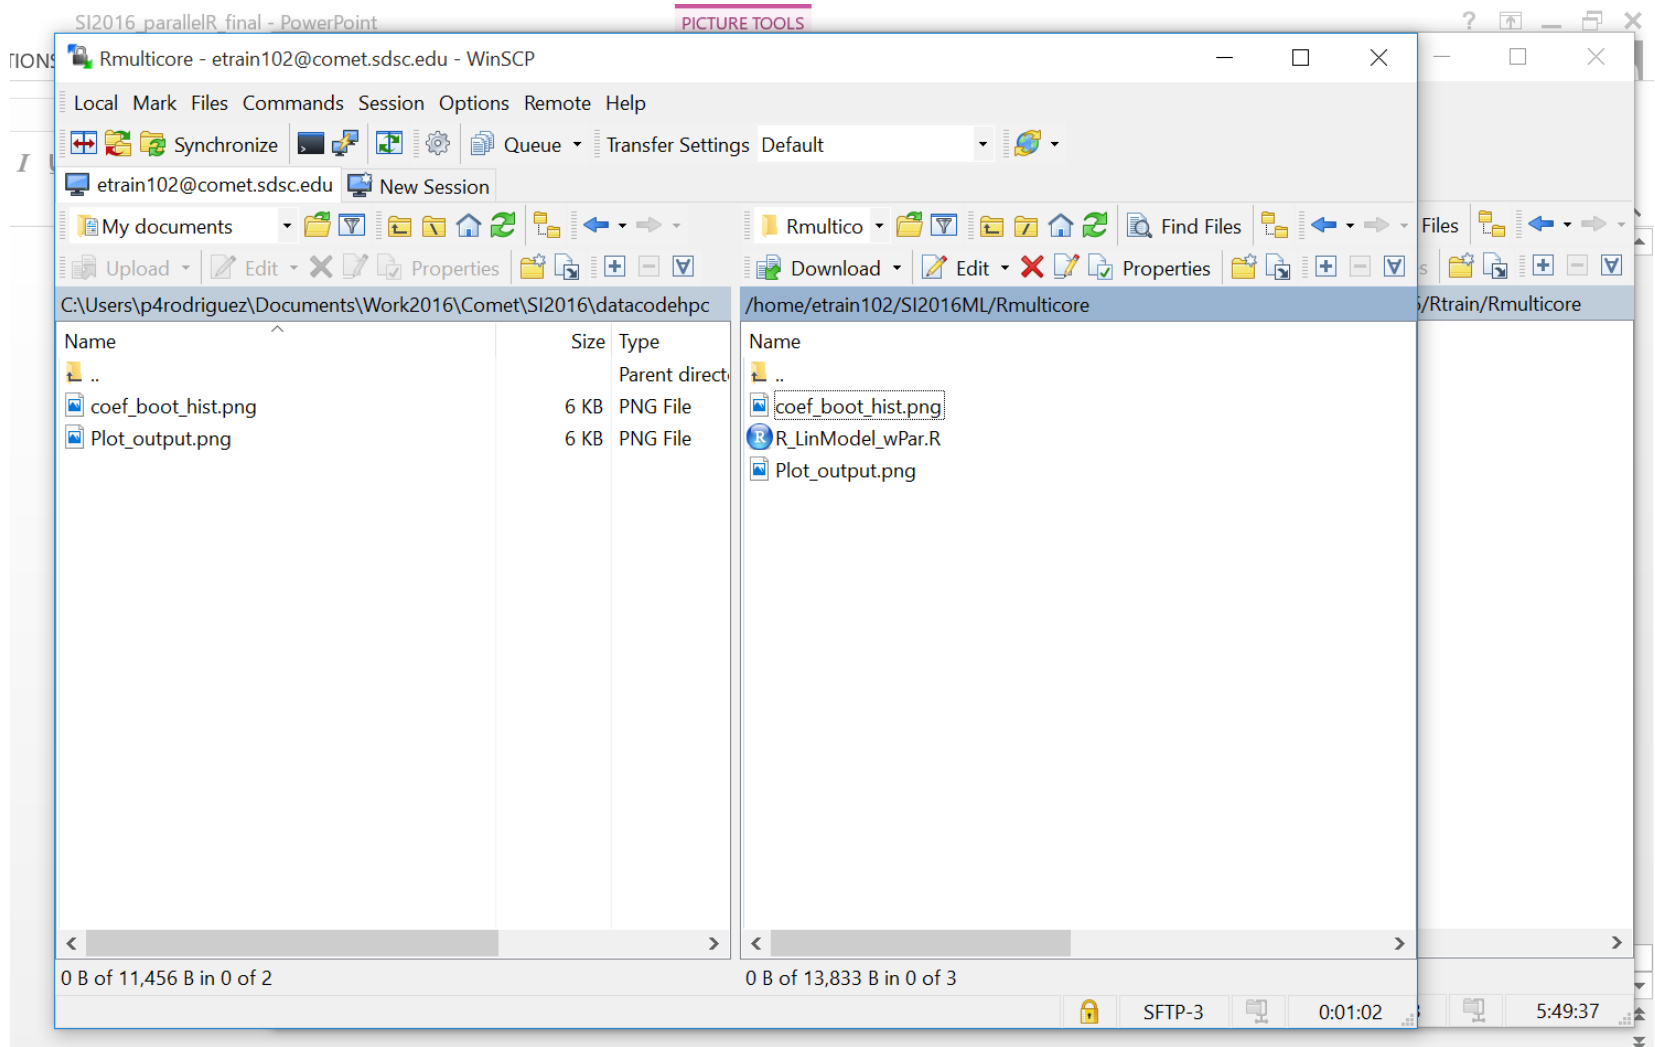
```
ptm <- proc.time()  
res_xboot = foreach(i = 1:num_resamp) %do% {  
  resamp_index = sample(1:N, size=sample_size, replace=T)  
  D_resamp     = D[resamp_index,] #pick out rows  
  lm_result    = lm(V1~.,data=D_resamp) #rerun model  
  coef_boot[i] = lm_result$coefficients['V10'] #save this beta  
}
```

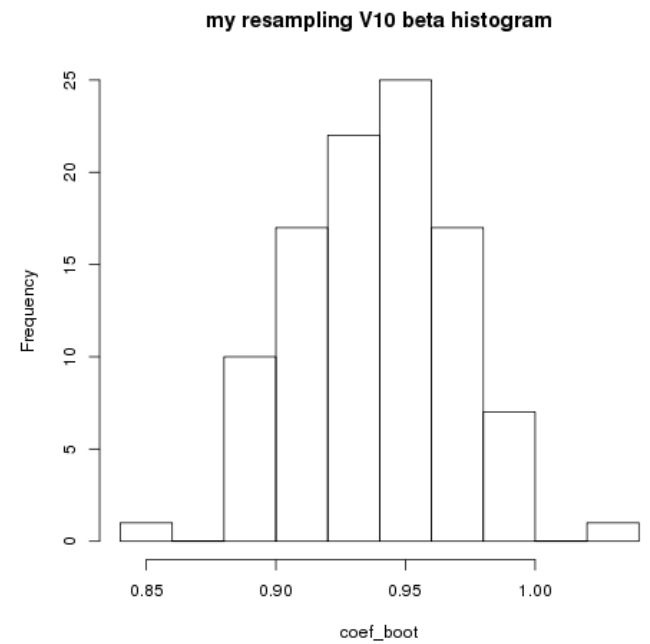
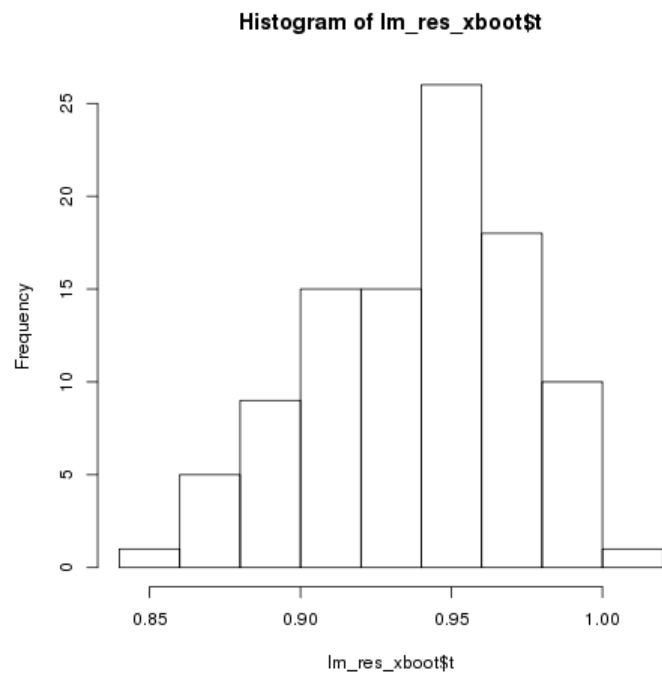
```
looptime=proc.time() - ptm #gather timing info  
print(paste('Time for do loop:',signif(looptime[3],5)))
```

[screen 0: bash] etrain102@comet-03-09:~/SI2016ML/Rmulticore

```
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> source('R_LinModel_wPar.R')
Loading required package: foreach
foreach: simple, scalable parallel programming from Revolution Analytics
Use Revolution R for scalability, fault tolerance and more.
http://www.revolutionanalytics.com
Loading required package: iterators
Loading required package: parallel
[1] "Time for do loop: 2.749"
> ls()
 [1] "B"           "coef_boot"    "D"            "D_resamp"     "i"
 [6] "lm_result"   "looptime"     "N"            "num_resamp"   "P"
[11] "ptm"         "res_xboot"    "resamp_index" "sample_size"  "X"
[16] "Y"
> str(coef_boot)
 num [1:100] 0.856 0.932 0.945 0.943 0.9 ...
> png('coef_boot_hist.png')
> hist(coef_boot,main='my resampling V10 beta histogram')
> dev.off()
null device
      1
>
```





Set to 1

And you
can
rerun and
compare
times

```
5 /oasis/scratch/comet/p4rodrig/temp_project/SI2016/Rtrain/Rmulticore/R_LinModel_wPar.R - p4rodrig@comet.sdsc.edu - E...
if (0) {  #if 0, then this section is blocked out,
-----
#Now try using the R Boot library
#you might need to install.packages('boot')
# parallel = c("no", "multicore", "snow") is the option we will play with
# -----
library(boot)
# 'boot' will call this and passes the data and indices
get_lm_res<-function(D_resamp,indices) {
  lm_resamp = lm(V1~.,data=D_resamp[indices,])  #rerun model
  return(lm_resamp$coefficient['V10'])
}
# -----
#call 'boot' function, a list of results from get_lm_res is returned
#the returned object has lots of boot information, $t is returned items from funtion
# -----
ptm =proc.time()
lm_res_xboot =boot(D, get_lm_res, R = num_resamp,parallel = "multicore")
looptime =proc.time() - ptm  #gather timing info
print(paste('Time for Bootstrap loop:',signif(looptime[3],5)))

#A sample of how to plot output to a file
png('Plot_output.png')
hist(lm_res_xboot$t)
dev.off()
# -----
Line: 74/76 Column: 767 Encoding: 1252 (ANSI - L
```