

Spark Machine Learning Library (mllib package)



MLLIB:

- Spark package for Machine Learning with resilient distributed datasets (RDD)
- Partly based on Python's 'scipy' package
- Some key ML algorithms (and growing)

```
[screen 0: bash] etrain102@comet-05-25:~/SI2016ML/SPML
Kill is control-H (^H).
[etrain102@comet-05-25 ~]$ ls
SI2016ML  vis
[etrain102@comet-05-25 ~]$ cd SI2016ML/SPML/
[etrain102@comet-05-25 SPML]$ source Setup_pyspark.sh
```

Enter:

```
> pyspark
```

after startup logs ...

Welcome to



version 1.3.0

```
Using Python version 2.6.6 (r266:84292, Feb 22 2013 00:00:18)
SparkContext available as sc, HiveContext available as sqlCtx.
```

```
In [1]:
```

Spark MLLIB Data Types

- **MLLIB works with RDD of:**
 - Arrays**
 - Vectors**
 - Labeled Points**

- Numpy package: Arrays

```
import numpy as np
```

```
x = np.array([1,2,3,4])
```

```
x[0]
```

```
Out[]: 1
```

- Array of arrays

```
x = np.array([[1,2],[3,4]])
```

```
x[0]
```

```
Out[]: array([1, 2])
```

A row

```
X[:,1]
```

```
Out[]: array([2,4])
```

A column

- MLLIB package: Vectors

```
from pyspark.mllib.linalg import Vectors
```

```
x = Vectors.dense([1,2,3,4])
```

```
x[0]
```

```
Out[]: 1
```

numpy arrays are interchangeable with mllib Vectors

- MLLIB package: RDD of Vectors

|

```
x = [ Vectors.dense([1,2,3,4]),
```

```
      Vectors.dense([5,6,7,8])]
```

```
xrdd = sc.parallelize(x)
```

```
xrdd
```

```
Out[]: <Python RDD .... >
```

now 'xrdd' has RDD actions available

- **MLLIB linalg package notes:**

SparseVectors also possible

**Distributed Matrix support in later pyspark releases
(but some available in Scala)**

- MLLIB package: LabeledPoint

```
from pyspark.mllib.regression import LabeledPoint
```

```
my_pt = LabeledPoint(1.0, Vectors.dense([1.0, 0.0, 3.0]))
```

```
my_pt.label
```

Class Label

```
Out[: 1.0
```

```
my_pt.features
```


```
Out[: [1.0, 0.0, 3.0]
```

Array

use this for setting up a class variable

RDD partitioning is along rows:

Large number of rows




...

...



Large number of Columns



⋮

⋮



Assume: data
partitioned on rows,
and 1 row fits in 1
computer memory

Spark MLLIB Clustering

- **MLLIB package: Kmeans**

Assign points to 'closest' cluster mean,
Update cluster mean,
Iterate until assignments converge

***Needs to iterate over data,
and calculate distance to cluster centers***

- Generating Random data:

```
from pyspark.mllib.random import RandomRDDs
```

normal distribution

```
c1_v=RandomRDDs.normalVectorRDD(
```

```
sc,20,2,
```

```
numPartitions=2,
```

20 rows, 2 columns

```
seed=1L).
```

```
map(lambda v:np.add([1,5],v))
```

center points around [1,5] by adding [1,5] to each point

- Generating Random data:

```
print c1_v.stats()
```

Ask for stats of the RandomRDD

```
Out[]: (count: 20, mean: [ 1.15426378  4.90223615],  
stdev: [ 1.10801385  1.40049822],  
max: [ 3.01083638  8.46783831],  
min: [-1.08338413  2.83934928])
```

You get basic stats
by column

- Generate 2 more classes and concatenate:

```
c2_v=RandomRDDs .... np.add([5,1],v))
```

```
c3_v=RandomRDDs ... np.add([4,6],v))
```

```
c12 =c1_v.union(c2_v)
```

```
my_data=c12.union(c3_v)
```

- MLLIB package:Kmeans

```
from pyspark.mllib.cluster import Kmeans, Kmeansmodel
```

- MLLIB package: Kmeans

```
my_kmmodel = KMeans.train(my_data,  
    k=3,  
    maxIterations=10,  
    runs=2,  
    initializationMode='k-means||')
```

number of clusters



use k-means over small, sample
to initialize (other option is 'random')



- Kmeans model functions:

```
#Sum Square Error of points to their cluster's center  
my_kmmodel.computeCost(my_data)
```

```
# get cluster centers  
my_kmmodel.clusterCenters
```

```
Out: [array([ 5.0476959 ,  1.27729209]), array([ 3.99839705,  
6.28073879]), array([ 0.95767935,  4.69770646])]
```

Note: with big data you sometimes only keep the cluster centers for further analysis

Spark MLLIB Classification

- **MLLIB package: Decision Tree**



Decision Tree induction:

At each node, partition data into bins
based on attribute values

- MLLIB package: Decision Tree



Decision Tree induction:

At each node, partition data into bins
based on attribute values

*Needs to iterate over data,
collect metrics,
choose nodes,
update current tree across*

weather data example

```
import numpy as np
from pyspark.mllib.linalg import Vectors
from pyspark.mllib.regression import LabeledPoint
```

Import modules

```
rawdata=[
    ['sunny',85,85,'FALSE',0],
    ['sunny',80,90,'TRUE',0],
    ['overcast',83,86,'FALSE',1],
    ['rainy',70,96,'FALSE',1],....
```

Raw data

```
data_df=sqlContext.createDataFrame(rawdata,  
                                   ['outlook','temp','humid','windy','play'])
```

```
#make RDD of labeled vectors (ie using only numbers!)
```

```
# build a dictionary to map outlook to new values
```

```
out2index={'sunny':0,'overcast':1,'rainy':2}
```

```
def newrow(dfrow):
```

```
    outnum = out2index.get(dfrow[0])
```

```
    outrow=list([outnum])
```

```
    outrow.append(dfrow[1]) #temp
```

```
    ....
```

```
    return (LabeledPoint(dfrow[4],outrow))
```

```
datax_rdd=data_df.map(newrow)
```

Convert categorical data (e.g. outlook=1,2 or 3) and make labeled points

- MLLIB package:DecisionTree

```
from pyspark.mllib.classification import  
DecisionTree
```

- MLLIB package: DecisionTree

```
from pyspark.mllib.classification import DecisionTree
```

```
my_dt = DecisionTree.trainClassifier(data_and_label,
```

```
2,
```

```
{0:3,3:2},
```

```
impurity='entropy',
```

```
maxDepth=5,
```

```
maxBins=32,
```

```
minInstancesPerNode=2)
```

2 classes,
Variable: Number of Categories

controls tree size and splitting
(these
may need cross validation to
optimize)

Confusion Matrix:

```
predictions = dt_model.predict(datax_rdd.map(lambda  
x: x.features))
```

```
labelsAndPredictions = datax_rdd.map(lambda lp:  
lp.label).zip(predictions)
```

```
Confusion_mat= [[ 5.  0.]  
[ 2.  7.]]
```

12 of 14 correct

Decision Tree output:

```
print dt_model.toDebugString()
```

IF-THEN-ELSE
rules are nodes

Out[]: DecisionTreeModel classifier of depth 3 with 9 nodes

If (feature 2 <= 80.0)

feat 2 is humid,
feat 1 is 'temp'

If (feature 1 <= 65.0)

Predict: 0.0

leaf node is
prediction

Else (feature 1 > 65.0)

Predict: 1.0

Else (feature 2 > 80.0)

If (feature 0 in {0.0})

pause