

Parallel Computing Using MPI & OpenMP

**Pietro Cicotti
(pcicotti@sdsc.edu)**

08/04/2015

Why Parallel Computing

- **Why we need high performance computing?**
 - Science!
 - Qualitative improvement in simulation resolution
 - Explosion of data to be analyzed
 - Both in industry and academic community
 - Competitive advantage
 - National Security
- **Why is parallel computing necessary?**



What is parallel computing?

- **Executing instructions concurrently on physical resources (not time slicing)**
 - Multiple tightly coupled resources (e.g. cores) collaboratively solving a single problem
- **Benefits**
 - Capacity
 - Memory, storage
 - Performance
 - More instructions per unit of time (FLOPS)
- **Cost and Complexity**
 - Coordinate tasks and resources
 - Use resources efficiently

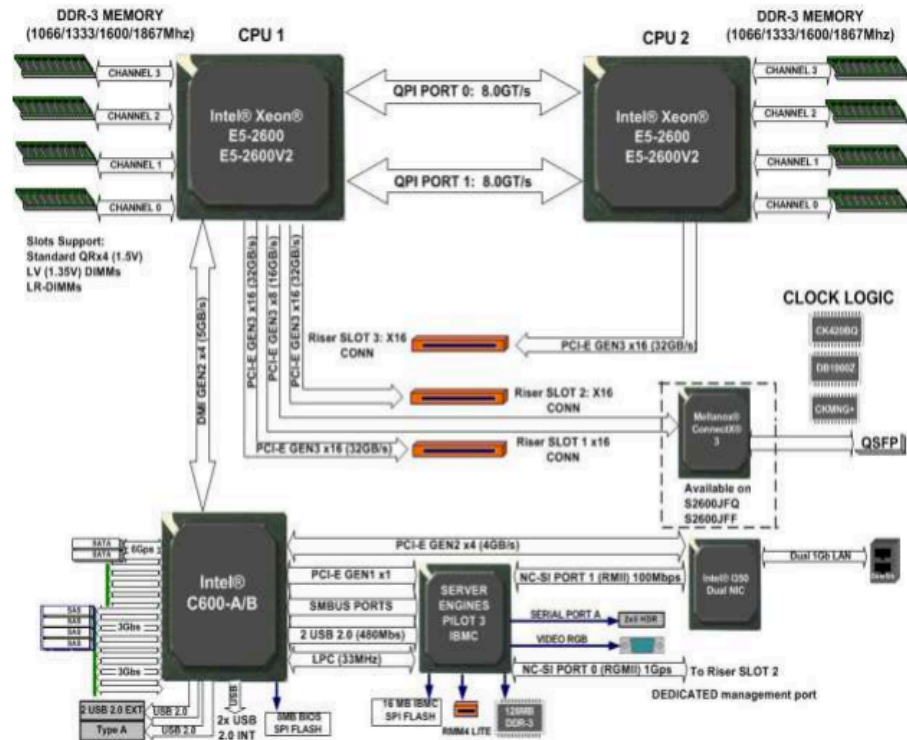
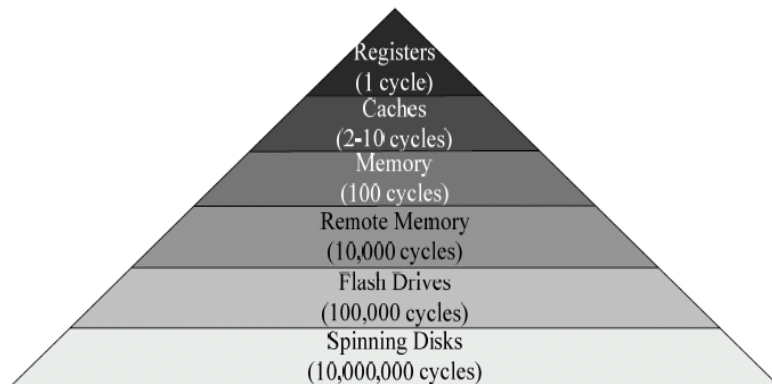
Flynn's Taxonomy

Single Instruction Single Data	Single Instruction Multiple Data
Multiple Instructions Single Data	Multiple Instructions Multiple Data

- **Single Program Multiple Data (SPMD)**
- **Multiple Program Multiple Data (MPMD)**


Today's clusters

- **Multi-socket server nodes**
 - NUMA, accelerators
- **High performance interconnect**
 - E.g. IB
- **Parallel File System and Storage System**
 - E.g. Lustre, GPFS



Distributed Memory Clusters Topologies

- **Mesh, Torus, Hypercube**
- **Tree based**
 - Fat-tree
 - Clos
- **Dragonfly**
- **Metrics**
 - Bandwidth
 - Diameter, Connectivity
 - Bisection bandwidth
- **Example**
 - Gordon
 - dual-rail
 - 3D torus of switches



Architecture, Memory, Models

- **Shared**
 - Communication model: share memory
- **Distributed**
 - Communication model: exchange messages
- **Execution**
 - Fork-Join (e.g. Thread Level Parallelism)
 - SPMD
- **Parallelism enabled by decomposing work**
 - Tasks can be executed concurrently
 - Some tasks can have dependencies

What is Multi-Threading?

- **A thread is a lightweight process**
 - OS entity
 - Shares virtual address space within the process
 - Share other resources (e.g. file descriptors, buffers)
 - Does not share a *context*
- **Fork and join are faster than for processes**
- **Fork and join are not free!**
- **Hardware threads support OS threads**
 - Hyperthreading

What is OpenMP

- **High level parallelism abstraction based on threads**
 - Easy to use
 - Suitable to an incremental approach
- **A specification**
 - “a portable, scalable model ... for developing portable parallel programs”
 - <http://openmp.org>
 - GNU, Intel, etc.
- **A set of**
 - Compiler directives
 - Library routines
 - Environment variables
- **Supports C/C++ and Fortran**

OpenMP Models

- **Fork/Join Execution**

- Process starts single threaded (master thread)
- Forks child threads activated in parallel regions (team)
- The team synchronizes and threads are disbanded
 - Overhead is mitigated by reusing threads
- Master thread continues execution of serial phases

- **Work decomposition**

- Explicit constructs
- Declarative in loops
 - Can be static or dynamic
 - Barriers and synchronization automatically inserted

Directives

- **Compiler directives apply to the succeeding structured block**
 - `#pragma omp`
 - Single statement or compound statement `{ }`
 - Clauses modify the properties of the directive
 - Compiler generate code
 - Instructions, functions, function calls
 - Transparent to the user
- **Main mechanism for declaring parallel regions of execution**
 - E.g. loops, sections

Regions, Loops, Sections, etc.

```
#pragma omp parallel [clause[ [, ]clause] ...] new-line  
structured-block  
clause:  
if(scalar-expression)  
num_threads(integer-expression)  
default(shared | none)  
private(list)  
firstprivate(list)  
shared(list)  
copyin(list)  
reduction(operator: list)
```

- **#pragma omp single/master**
- **simd**
- **tasks**

```
#pragma omp for [clause[[,] clause] ... ] new-line for-loops  
clause:  
private(list)  
firstprivate(list)  
lastprivate(list)  
reduction(operator: list)  
schedule(kind[, chunk_size])  
collapse(n)  
ordered  
nowait
```

```
#pragma omp sections [clause[[,] clause] ...] new-line  
{  
  #pragma omp section  
  structured-block  
  ...  
}  
clause:  
private(list)  
firstprivate(list)  
lastprivate(list)  
reduction(operator: list)  
nowait
```

Scope of Variables

- **Clauses determine the scope of variables**
 - Default: shared (external)
- **private**
 - Also if declared inside region
- **firstprivate**
- **shared**
- **lastprivate**
- **reductions**
- **default**

Decomposition of the Iteration Space

schedule(kind[,chunk_size])

kind:

static: Iterations are divided into chunks of size *chunk_size* and assigned to threads in the team in round-robin fashion in order of thread number.

dynamic: Each thread executes a chunk of iterations then requests another chunk until none remain.

guided: Each thread executes a chunk of iterations then requests another chunk until no chunks remain to be assigned.

auto: The decision regarding scheduling is delegated to the compiler and/or run me system.

runtime: The schedule and chunk size are taken from the *run-sched-var* ICV.

More Synchronization

- **#pragma omp critical**
 - Executed by one thread at a time
- **#pragma omp barrier**
 - Explicit barrier
- **#pragma omp atomic**
 - Atomic instruction
 - Storage is accessed atomically

Controlling and Querying the environment

- **Env vars**
 - e.g. OMP_NUM_THREADS
- **Routines**
 - Execution
 - omp_[get | set]_num_threads
 - omp_get_thread_num
 - Locking
 - omp_init_lock, omp_set_lock, omp_unset_lock
 - Timing
 - omp_get_wtime()
 - omp_get_wtick

Compute PI with openmp

- **Examples on Gordon**
 - /home/diag/opt/Sl2016/openmp/
 - qsub -l -lnodes=1:ppn=16:native,walltime=00:60:00

1D heat equation with OpenMP

Correctness Considerations

- **Sharing data**

- Dependencies must be enforced
- Operations are not atomic unless specified

Thread 0	Thread 1
x=0	
++X	++X
x==2?	

- Caches are coherent, registers are not
- **Loops may carry dependencies across iterations**
 - for i=0 to 9 do A[i]+=B[i]
 - for i=0 to 9 do A[i]+=A[i+1] (try with a unit array)

Performance Considerations

- **Synchronizations and serialization hurt performance**
 - barriers, locks, critical sections, single thread blocks
 - nowait close
- **Coarse parallelization reduces overhead**
- **Preserve locality**
 - NUMA
 - Bind threads to cores
 - Avoid false sharing
- **Use optimal scheduling**

False Sharing

- **Modern processors have SRAM caches**
 - Low capacity
 - High performance
 - e.g. 1 cycle, 20 cycle, 100 cycles, 300 cycles
- **Caches**
 - size, associativity, line (IA64 uses 64B cache lines)
 - `x[0]+=1; x[7]+=1; // may be same line`
- **Shared memory**
 - Coherency preserved by coherency protocol
 - Invalidate copies when writing
 - Writes cause coherency traffic and serialization
- **Severe impact on performance!**

False sharing: example

```
#pragma omp parallel for schedule(static,1)  
for(int i=0; i<N; ++i)  
    ++x[i];
```

- **Solutions**

- align data and partition boundaries to cache line size
 - `int x __attribute__((aligned (16))) = 0;`
- pad arrays when needed
 - Element are cache line aligned
 - Boundaries are cache line aligned
- Use local copies whenever possible

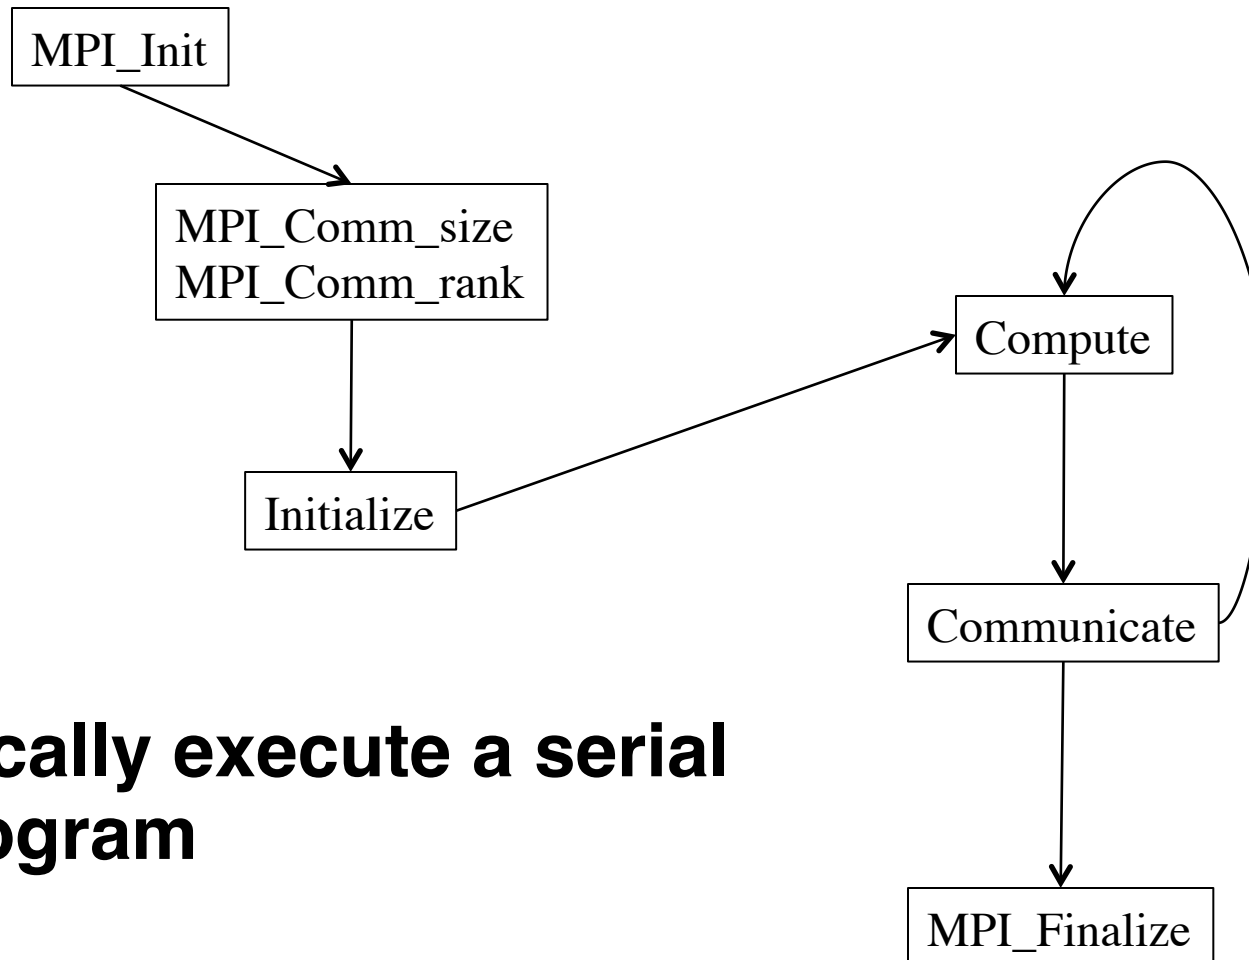
Practice!

- **Check the specification**
 - <http://www.openmp.org>
- **Try to write a program to do a parallel sort**
 - d&c: quick sort (man qsort), serial merge
 - Can you improve on serial merge?
 - Amdhal's law?
- **Code in /home/diag/opt/SI2016**

Message Passing Interface (MPI)

- **Low level message passing abstraction**
 - SPMD execution model + messages
 - Designed for distributed memory
 - send-recv basic primitives
- **MPI: API specification**
 - Portable: de-fact standard for parallel computing
 - <http://www.mpi-forum.org>
 - E.g. openMPI, mpich, mvapich, LAM
 - High performance implementations available virtually on any interconnect and system
 - Point-to-point communication, datatypes, collective operations
 - One-sided communication, Parallel file I/O, Tool support, ...

Bulk Synchronous Programming with MPI



- **Locally execute a serial program**

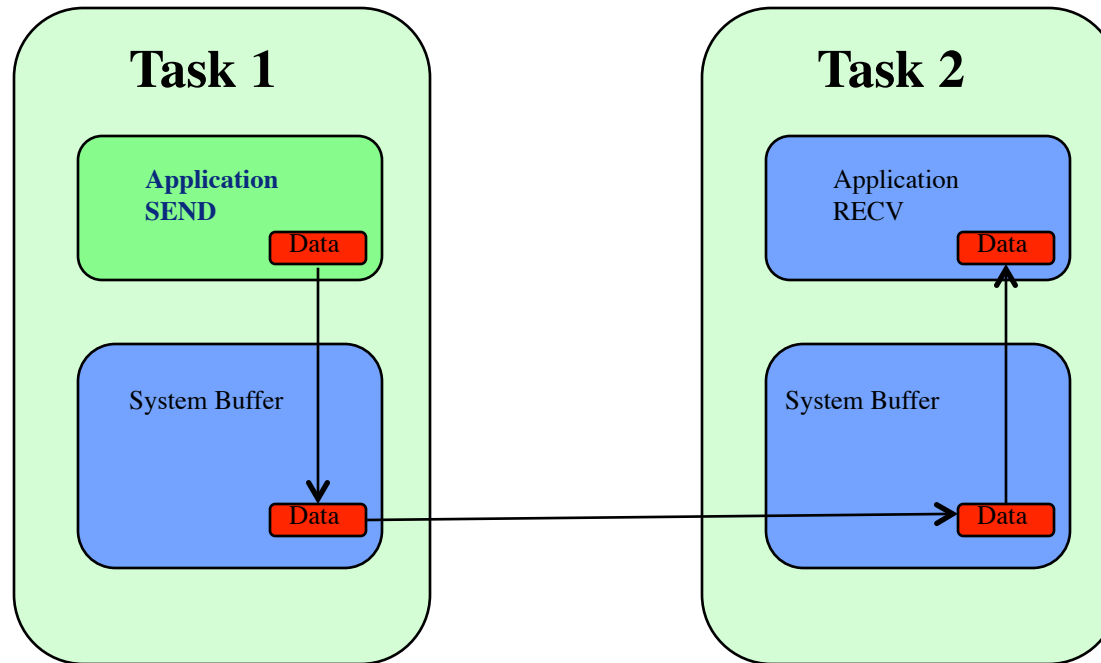
Communicators

- **Define a communication domain**
 - set of processes that communicate with each other
 - Required for message transfer routines
- **MPI_COMM_WORLD**
 - Default communicator
 - Includes all the processes
- **Useful for library developers**
- **Logically partition the data/processes**
 - Match data and work decomposition
- **MPI_Comm_size, MPI_Comm_rank**

Point-to-Point Communication

- **MPI_[I][?]Send, MPI_[I]Recv**
- **Message = data + envelop (src,dst,comm,tag)**
- **[?] communication mode modifies the semantics of the send**
 - Standard
 - Buffered
 - Synchronous
 - Ready
- **[I] Immediate routines**
 - Blocking vs non-blocking
 - Start, wait, test
- **Example**
 - `MPI_Send(buf, 1, MPI_INT, 1, 0, MPI_COMM_WORLD)`
 - `MPI_Recv(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status)`

Buffering and message transfer



Blocking Send-Recv

Avoiding deadlocks

- **Deadlocks are common mistakes**
 - Unexpected behavior/semantics
 - Circular dependencies
- **Example (try different modes!):**

```
if(myrank) {  
    MPI_Ssend(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);  
    MPI_Recv(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);  
} else {  
    MPI_Ssend(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);  
    MPI_Recv(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);  
}
```

Fixes to deadlock example?

```
if(myrank) {  
    MPI_Ssend(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);  
    MPI_Recv(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);  
} else {  
    MPI_Recv(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);  
    MPI_Ssend(buf, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);  
}
```

```
MPI_Send(buf, 1, MPI_INT, myrank?0:1, 0, MPI_COMM_WORLD);  
MPI_Recv(buf, 1, MPI_INT, myrank?0:1, 0, MPI_COMM_WORLD, &status);
```

```
MPI_Isend(buf, 1, MPI_INT, myrank?0:1, 0, MPI_COMM_WORLD, &req);  
MPI_Recv(buf, 1, MPI_INT, myrank?0:1, 0, MPI_COMM_WORLD, &status);  
MPI_Wait(&req, &status);
```


Compute PI with MPI

Collective Communication


- **All ranks in a communicator participate**
 - Potential optimizations with respect to point-to-point
 - Broadcast: $n-1$ messages vs. $\log(n)$ messages
- **Barriers**
 - Synchronize all ranks
- **Broadcast**
- **Reduction**
- **Gather/Scatter, Alltoall**
- **Scan**
- ***All* and *vector* variants**

1D heat equation with MPI

- $\partial T / \partial t = \alpha (\partial^2 T / \partial x^2)$; $T(0) = 0$; $T(1) = 0$; ($0 \leq x \leq 1$)
- $T(x, 0)$ initial condition
- Discretization
 - $T(x_i, n+1) - T(x_i, n) = (\alpha \Delta t / \Delta x^2) (T(x_{i-1}, n) - 2T(x_i, n) + T(x_{i+1}, n))$
- Partitioning
 - Ghost cells



Simple Application using MPI: 1-D Heat Equation



Processor 0:

Local Data Index : $ilocal = 0, 1, 2, 3, 4$

Global Data Index: $iglobal = 0, 1, 2, 3, 4$

Solve the equation at (1,2,3)

Data Exchange: Get 4 from processor 1; Send 3 to processor 1

Processor 1:

Local Data Index : $ilocal = 0, 1, 2, 3, 4$

Global Data Index : $iglobal = 3, 4, 5, 6, 7$

Solve the equation at (4,5,6)

Data Exchange: Get 3 from processor 0; Get 7 from processor 2; Send 4 to processor 0; Send 6 to processor 2

Processor 2:


Local Data Index : $ilocal = 0, 1, 2, 3, 4$

Global Data Index : $iglobal = 6, 7, 8, 9, 10$

Solve the equation at (7,8,9)

Data Exchange: Get 6 from processor 1; Send 7 to processor 1

Simple Application using MPI: 1-D Heat Equation



% more data0.dat

Processor 0

```
ilocal= 0 ;iglobal= 0 ;T= 0.000000000000000000E+00
ilocal= 1 ;iglobal= 1 ;T= 0.307205621017284991
ilocal= 2 ;iglobal= 2 ;T= 0.584339815421976549
ilocal= 3 ;iglobal= 3 ;T= 0.804274757358271253
ilocal= 4 ;iglobal= 4 ;T= 0.945481682332597884
```

% more data2.dat

Processor 2

```
ilocal= 0 ;iglobal= 6 ;T= 0.945481682332597995
ilocal= 1 ;iglobal= 7 ;T= 0.804274757358271253
ilocal= 2 ;iglobal= 8 ;T= 0.584339815421976660
ilocal= 3 ;iglobal= 9 ;T= 0.307205621017285102
ilocal= 4 ;iglobal= 10 ;T= 0.000000000000000000E+00
```


% more data1.dat

Processor 1

```
ilocal= 0 ;iglobal= 3 ;T= 0.804274757358271253
ilocal= 1 ;iglobal= 4 ;T= 0.945481682332597884
ilocal= 2 ;iglobal= 5 ;T= 0.994138272681972301
ilocal= 3 ;iglobal= 6 ;T= 0.945481682332597995
ilocal= 4 ;iglobal= 7 ;T= 0.804274757358271253
```

Fortran MPI Code: 1-D Heat Equation

Simple Application using MPI: 1-D Heat Equation



- Compilation

Fortran: `mpif90 -nofree -o heat_mpi.exe heat_mpi.f90`

- Run Job:

`qsub heat_mpi.cmd`

Performance Considerations

- **Overlap communication with computation**
 - Use non-blocking primitives
 - Hide communication cost
 - Split-phase programming
- **Minimize surface-to-volume ratio**
 - Ghost cell exchange
- **Avoid communication**
 - Even at the cost of some more computation
 - Example: double size of ghost cell and communicate every other time step

Debugging, Profiling, Tracing

- **Hard to use command line debuggers**
 - gdb, idb
- **Allinea DDT is installed on Gordon and Comet**
 - GUI
- **Profiling**
 - mpiP, TAU, IPM installed on Gordon and Comet
- **Useful information**
 - runtime breakdown: communication vs. computation
- **Identify bottlenecks and scaling issues**

mpiP sample output

Compile

```
mpif90 -nofree -g -o heat_mpi_profile.exe heat_mpi.f90 -L/home/diag/opt/mpiP/v3.4.1/lib -lmpiP -L/opt/gnu/lib -lbfd -lz -liberty
```

Run

```
mpirun_rsh -hostfile $PBS_NODEFILE -np 3 ./heat_mpi_profile.exe
```

```
@ mpiP
@ Command : ./heat_mpi_profile.exe
@ Version      : 3.4.1
@ MPIP Build date   : Aug 3 2014, 19:18:28
@ Start time       : 2014 08 06 08:50:44
@ Stop time        : 2014 08 06 08:50:44
@ Timer Used       : PMPI_Wtime
@ MPIP env var      : [null]
@ Collector Rank    : 0
@ Collector PID     : 53941
@ Final Output Dir  : .
@ Report generation : Single collector task
@ MPI Task Assignment : 0 gcn-13-35.sdsc.edu
@ MPI Task Assignment : 1 gcn-13-35.sdsc.edu
@ MPI Task Assignment : 2 gcn-13-35.sdsc.edu
```

mpiP Output

@--- MPI Time (seconds) -----

Task	AppTime	MPITime	MPI%
0	0.0702	0.00513	7.30
1	0.0728	0.00516	7.09
2	0.0732	0.00519	7.08
*	0.216	0.0155	7.16

@--- Callsites: 1 -----

ID	Lev	File/Address	Line	Parent_Funct	MPI_Call
1	0	0x40dbf4		main	Send

@--- Aggregate Time (top twenty, descending, milliseconds) -----

Call	Site	Time	App%	MPI%	COV
Send	1	15.2	7.04	98.34	0.00
Recv	1	0.257	0.12	1.66	0.23

mpiP output

@--- Aggregate Sent Message Size (top twenty, descending, bytes) -----

Call	Site	Count	Total	Avrg Sent%
Send	1	12	96	8 100.00

@--- Callsite Time statistics (all, milliseconds): 6 -----

Name	Site	Rank	Count	Max	Mean	Min	App%	MPI%
Recv	1	0	3	0.052	0.0233	0.008	0.10	1.37
Recv	1	1	6	0.026	0.0132	0.004	0.11	1.53
Recv	1	2	3	0.057	0.036	0.02	0.15	2.08
Send	1	0	3	5.05	1.69	0.004	7.20	98.63
Send	1	1	6	5.06	0.847	0.003	6.98	98.47
Send	1	2	3	5.07	1.69	0.004	6.93	97.92
Send	1	*	24	5.07	0.645	0.003	7.16	100.00

@--- Callsite Message Sent statistics (all, sent bytes) -----

Name	Site	Rank	Count	Max	Mean	Min	Sum
Send	1	0	3	8	8	8	24
Send	1	1	6	8	8	8	48
Send	1	2	3	8	8	8	24
Send	1	*	12	8	8	8	96

@--- End of Report -----

Data Types

C Data Types		FORTRAN Data Types
MPI_CHAR	MPI_C_DOUBLE_COMPLEX	MPI_CHARACTER
MPI_WCHAR	MPI_C_LONG_DOUBLE_COMPLEX	MPI_INTEGER
MPI_SHORT	MPI_C_BOOL	MPI_INTEGER1
MPI_INT	MPI_LOGICAL	MPI_INTEGER2
MPI_LONG	MPI_C_LONG_DOUBLE_COMPLEX	MPI_INTEGER4
MPI_LONG_LONG_INT	MPI_INT8_T	MPI_REAL
MPI_LONG_LONG	MPI_INT16_T	MPI_REAL2
MPI_SIGNED_CHAR	MPI_INT32_T	MPI_REAL4
MPI_UNSIGNED_CHAR	MPI_INT64_T	MPI_REAL8
MPI_UNSIGNED_SHORT	MPI_UINT8_T	MPI_DOUBLE_PRECISION
MPI_UNSIGNED_LONG	MPI_UINT16_T	MPI_COMPLEX
MPI_UNSIGNED	MPI_UINT32_T	MPI_DOUBLE_COMPLEX
MPI_FLOAT	MPI_UINT64_T	MPI_LOGICAL
MPI_DOUBLE	MPI_BYTE	MPI_BYTE
MPI_LONG_DOUBLE	MPI_PACKED	MPI_PACKED
MPI_C_COMPLEX		
MPI_C_FLOAT_COMPLEX		

MPI Reduction Operations

NAME	OPERATION
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI_LAND	Logical AND
MPI_BAND	Bit-wise AND
MPI_LOR	Logical OR
MPI_BOR	Bit-wise OR
MPI_LXOR	Logical XOR
MPI_BXOR	Bit-wise XOR
MPI_MAXLOC	Maximum value and location
MPI_MINLOC	Minimum value and location

Advanced MPI

- **One-sided communication**
- **Derived data types**
- **Parallel I/O**
- **Groups, topologies, and communicators management**
- **Dynamic process creation and management**
- **Tools support**

Practice

- **Code in /home/diag/opt/Sl2016**
- **Find bugs in code sample.f**
 - Compile mpif90 -o sample sample.f
- **Questions?**
- **Try your favorite computation**

References

- **Excellent tutorials from LLNL:**
 - <https://computing.llnl.gov/tutorials/mpi/>
 - <https://computing.llnl.gov/tutorials/openMP/>
- **MPI for Python:**
 - <http://mpi4py.scipy.org/docs/usrman/tutorial.html>
- **MVAPICH2 User Guide:**
 - <http://mvapich.cse.ohio-state.edu/userguide/>